



## **Assignment-2**

Joseline Luvena Tanujaya

Sweta Bansal

Vibhanshu



# IDS 572 - Assignment 2

Joseline Tanujaya, Sweta Bansal, Vibhanshu

```
#split the data into trn, tst subsets
set.seed(123)
nr=nrow(lcdf)
trnIndex = sample(1:nr, size = round(0.7*nr), replace=FALSE)
lcdfTrn=lcdf[trnIndex,]
lcdfTst = lcdf[-trnIndex,]

dim(lcdfTrn)
```

```
## [1] 56715    46
```

```
dim(lcdfTst)
```

```
## [1] 24307    46
```

## Predicting Loan Status with GLM

In building the glm model, we adjusted several parameters: type.measure, alpha, and weights (balanced data).

Below is the complete tabulation of the parameters and their resulting AUC.

##	X	X.1	X.2	Lambda.min	X.3	Lambda.1se	X.4
## 1	type.measure	alpha	balanced	AUC on Trn	AUC on Tst	AUC on Trn	AUC on Tst
## 2	deviance	1	N	0.6954079	0.6954079	0.687354	0.687354
## 3	auc	1	N	0.6954079	0.6910984		
## 4	class	1	N	0.6854198	0.6809635		
## 5	deviance	0	N	0.6946417	0.6908809		
## 6	auc	0	N	0.6946417	0.6908809		
## 7	class	0	N	0.692251	0.6899955		
## 8	deviance	1	Y	0.6954781	0.6910646		
## 9	deviance	0.3	Y	0.6954573	0.6910745		
## 10	deviance	0.5	Y	0.6954791	0.6910773		
## 11	deviance	0.7	Y	0.6954865	0.6910773		

We used AUC measures to define the best glm model. The models that we got are not too different with each other. However with very slight differences, we can still determine the best model. The best model is the one with balanced class of loan status (using weights), with **type.measure = "deviance"**, **alpha = 1** with AUC of 0.6953326.

We ran the code several times, and the model with balanced weights always give slightly better results. However, depending on the Train and Test data we build and test the models on, alpha = 1 or alpha = 0.7 can be the best options. When comparing the glm model with other models (ranger, xgboost), we will use alpha = 1.

For variable selection method, we will scale the coefficients using the Lasso and Ridge regularization parameters. This is done by setting different values to alpha. When alpha = 1, we are using purely Lasso (L1). When alpha = 0, we are using purely Ridge (L2). Any value between 0-1 is attributed to the percentage using L1. For example, if

alpha = 0.2, then  $L1 = 0.2$  and  $L2 = (1 - 0.2) = 0.8$ . We have tried several values between 0-1, but the best result still comes from alpha = 1 (Lasso only). Lasso puts pressure on the coefficients to approach 0, therefore we are discarding some variables with low importance.

We experimented using lambda 1se using the best models (with balanced weights). Lambda min has proven to be the better lambda option, since it has higher accuracy. There is no overfitting problems with our models, therefore we do not need to use a bigger lambda than lambda min.

The experimentation with type.measure will be documented at section 1B.

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.0.4
```

```
## Loading required package: Matrix
```

```
##  
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':  
##  
##   expand, pack, unpack
```

```
## Loaded glmnet 4.1-1
```

```
library(tidyverse)  
library(lubridate)  
library(tidyr)  
library(caret)
```

```
## Loading required package: lattice
```

```
##  
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':  
##  
##   lift
```

```
library(prediction)
```

```
## Warning: package 'prediction' was built under R version 4.0.4
```

```
library(ROCR)
```

```
##  
## Attaching package: 'ROCR'
```

```
## The following object is masked from 'package:prediction':  
##  
## prediction
```

```
library(rpart)  
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##  
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':  
##  
## cov, smooth, var
```

```
library(broom)  
  
levels(lcdfTrn$loan_status)
```

```
## [1] "Fully Paid" "Charged Off"
```

```
yTrn<-factor(if_else(lcdfTrn$loan_status=="Fully Paid", '1', '0') )  
yTst<-factor(if_else(lcdfTst$loan_status=="Fully Paid", '1', '0') )  
  
xDTTrn<-lcdfTrn %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_pymnt)  
xDTst<-lcdfTst %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_pymnt)  
  
#Use Lasso regularization (alpha = 1, default)  
glm1s_cv<- cv.glmnet(data.matrix(xDTTrn), yTrn, family="binomial")  
  
glm1s_cv$lambda.min
```

```
## [1] 0.0001133834
```

```
glm1s_cv$lambda.1se
```


```
## [1] 0.006193316
```

```
#get the variables with non-zero coefficients from the regularized model  
nzCoef<-tidy(coef(glm1s_cv, s= glm1s_cv$lambda.1se))
```

```
## Warning: 'tidy.dgCMatrix' is deprecated.  
## See help("Deprecated")
```

```
## Warning: 'tidy.dgTMatrix' is deprecated.  
## See help("Deprecated")
```

```
nzCoefVars <- nzCoef[-1,1]  
  
plot(glm1s_cv)
```



```
coef(glm1s_cv, s = glm1s_cv$lambda.min)
```

```
## 42 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  3.462405e+00
## loan_amnt    .
## int_rate     -3.618328e-02
## installment  -2.326383e-04
## grade        -7.475791e-02
## sub_grade     -3.770364e-02
## emp_length    1.195171e-02
## home_ownership -6.004841e-02
## annual_inc    -3.735868e-07
## purpose       3.193568e-03
## dti           -1.840644e-02
## initial_list_status -2.126423e-02
## total_rev_hi_lim  4.428945e-06
## acc_open_past_24mths -4.566196e-02
## avg_cur_bal    1.358973e-06
## bc_open_to_buy -4.258079e-06
## bc_util        -2.839648e-03
## chargeoff_within_12_mths -3.864755e-02
## delinq_amnt    2.765696e-05
## mo_sin_old_il_acct -1.073440e-04
## mo_sin_old_rev_tl_op  2.288770e-04
## mo_sin_rcnt_rev_tl_op .
## mo_sin_rcnt_tl    2.409513e-03
## mort_acc        2.626384e-02
## mths_since_recent_bc -1.606128e-04
## mths_since_recent_inq  3.707635e-03
## num_bc_sats     .
## num_bc_tl      -3.728922e-02
## num_il_tl       5.091431e-03
## num_op_rev_tl   -2.973017e-02
## num_rev_accts   2.132643e-02
## num_sats        4.396394e-03
## num_tl_30dpd    -2.203274e-01
## pct_tl_nvr_dlq   1.377219e-03
## tax_liens       -2.946152e-02
## tot_hi_cred_lim  5.718190e-07
## total_bal_ex_mort -4.584606e-06
## total_bc_limit   7.652567e-06
## total_il_high_credit_limit  4.809257e-06
## propSatisBankcardAccts -1.311651e-01
## prop_OpAccts_to_TotAccts  3.519362e-01
## propLoanAmt_to_AnnInc -4.766838e-01
```

```
coef(glm1s_cv, s = glm1s_cv$lambda.1se)
```

```
## 42 x 1 sparse Matrix of class "dgCMatrix"
##                               1
## (Intercept)                3.360645e+00
## loan_amnt                   .
## int_rate                    -2.977258e-02
## installment                 .
## grade                       -4.364347e-02
## sub_grade                   -5.318270e-02
## emp_length                  .
## home_ownership              -3.442481e-02
## annual_inc                  .
## purpose                     .
## dti                         -1.009910e-02
## initial_list_status         .
## total_rev_hi_lim            .
## acc_open_past_24mths        -3.572049e-02
## avg_cur_bal                 .
## bc_open_to_buy              3.792271e-07
## bc_util                     -9.069142e-04
## chargeoff_within_12_mths    .
## delinq_amnt                 .
## mo_sin_old_il_acct          .
## mo_sin_old_rev_tl_op        .
## mo_sin_rcnt_rev_tl_op       .
## mo_sin_rcnt_tl              .
## mort_acc                    1.062411e-02
## mths_since_recent_bc        .
## mths_since_recent_inq       1.901960e-04
## num_bc_sats                  .
## num_bc_tl                   .
## num_il_tl                   .
## num_op_rev_tl               .
## num_rev_accts               .
## num_sats                     .
## num_tl_30dpd                .
## pct_tl_nvr_dlq              .
## tax_liens                   .
## tot_hi_cred_lim              6.093373e-07
## total_bal_ex_mort           .
## total_bc_limit              .
## total_il_high_credit_limit   .
## propSatisBankcardAccts      .
## prop_OpAccts_to_TotAccts     .
## propLoanAmt_to_AnnInc       -4.149634e-01
```

```
#find the index of the best lambda
which(glmnet.cv$lambda == glmnet.cv$lambda.1se)
```

```
## [1] 28
```

```
glmnet.cv$glmnet.fit$dev.ratio[which(glmnet.cv$lambda == glmnet.cv$lambda.1se) ]
```

```
## [1] 0.06481165
```


```
glm1s_cv$glmnet.fit
```




```
##
## Call: glmnet(x = data.matrix(xDTrn), y = yTrn, family = "binomial")
##
##      Df %Dev   Lambda
## 1    0 0.00 0.076350
## 2    1 0.95 0.069570
## 3    1 1.72 0.063390
## 4    1 2.36 0.057760
## 5    1 2.88 0.052630
## 6    1 3.32 0.047950
## 7    1 3.68 0.043690
## 8    1 3.97 0.039810
## 9    2 4.22 0.036270
## 10   2 4.42 0.033050
## 11   2 4.60 0.030120
## 12   2 4.74 0.027440
## 13   3 4.86 0.025000
## 14   3 4.95 0.022780
## 15   5 5.08 0.020760
## 16   6 5.27 0.018910
## 17   6 5.45 0.017230
## 18   6 5.61 0.015700
## 19   6 5.75 0.014310
## 20   8 5.87 0.013040
## 21   9 5.99 0.011880
## 22   9 6.09 0.010820
## 23   9 6.18 0.009862
## 24  10 6.25 0.008985
## 25  10 6.32 0.008187
## 26  10 6.38 0.007460
## 27  11 6.43 0.006797
## 28  12 6.48 0.006193
## 29  14 6.53 0.005643
## 30  16 6.58 0.005142
## 31  16 6.63 0.004685
## 32  16 6.67 0.004269
## 33  16 6.71 0.003890
## 34  17 6.74 0.003544
## 35  17 6.76 0.003229
## 36  17 6.79 0.002942
## 37  17 6.80 0.002681
## 38  17 6.82 0.002443
## 39  20 6.83 0.002226
## 40  20 6.85 0.002028
## 41  20 6.86 0.001848
## 42  22 6.87 0.001684
## 43  24 6.88 0.001534
## 44  26 6.89 0.001398
## 45  28 6.90 0.001274
## 46  30 6.91 0.001161
## 47  31 6.92 0.001057
## 48  31 6.93 0.000963
## 49  32 6.94 0.000878
## 50  33 6.95 0.000800
```

```
## 51 34 6.96 0.000729
## 52 35 6.98 0.000664
## 53 36 6.99 0.000605
## 54 36 7.00 0.000551
## 55 37 7.01 0.000502
## 56 37 7.02 0.000458
## 57 37 7.03 0.000417
## 58 38 7.03 0.000380
## 59 38 7.04 0.000346
## 60 38 7.04 0.000316
## 61 39 7.05 0.000288
## 62 39 7.05 0.000262
## 63 39 7.06 0.000239
## 64 39 7.06 0.000218
## 65 39 7.06 0.000198
## 66 38 7.07 0.000180
## 67 38 7.07 0.000164
## 68 38 7.07 0.000150
## 69 38 7.07 0.000137
## 70 38 7.07 0.000124
## 71 38 7.07 0.000113
```

```
plot(glm1s_cv$glmnet.fit)
```



```
plot(glm1s_cv$glmnet.fit, xvar="lambda")
```



```
#as.matrix(coef(glmLs_cv, s = glmLs_cv$lambda.min))
#as.matrix(coef(glmLs_cv, s = glmLs_cv$lambda.1se))
```

```
#the lambda values used are in glmLs_cv$lambda
glmLs_cv$lambda
```

```
## [1] 0.0763540922 0.0695710051 0.0633905088 0.0577590707 0.0526279140
## [6] 0.0479525952 0.0436926188 0.0398110870 0.0362743798 0.0330518638
## [11] 0.0301156273 0.0274402380 0.0250025229 0.0227813677 0.0207575338
## [16] 0.0189134917 0.0172332692 0.0157023131 0.0143073629 0.0130363362
## [21] 0.0118782240 0.0108229953 0.0098615102 0.0089854409 0.0081871991
## [26] 0.0074598709 0.0067971565 0.0061933158 0.0056431187 0.0051417996
## [31] 0.0046850163 0.0042688124 0.0038895828 0.0035440430 0.0032292000
## [36] 0.0029423268 0.0026809386 0.0024427714 0.0022257623 0.0020280317
## [41] 0.0018478670 0.0016837076 0.0015341316 0.0013978436 0.0012736630
## [46] 0.0011605143 0.0010574174 0.0009634794 0.0008778865 0.0007998975
## [51] 0.0007288368 0.0006640890 0.0006050931 0.0005513383 0.0005023589
## [56] 0.0004577308 0.0004170672 0.0003800161 0.0003462565 0.0003154961
## [61] 0.0002874683 0.0002619304 0.0002386612 0.0002174592 0.0001981407
## [66] 0.0001805384 0.0001644999 0.0001498862 0.0001365707 0.0001244381
## [71] 0.0001133834
```

```
# and the cross-validation 'loss' at each lambda is in glmLs_cv$cvm
glmLs_cv$cvm
```

```
## [1] 0.8304065 0.8228304 0.8163836 0.8110919 0.8067410 0.8031442 0.8001722
## [8] 0.7977096 0.7956683 0.7939745 0.7925666 0.7913910 0.7904136 0.7895797
## [15] 0.7885381 0.7870839 0.7855542 0.7842445 0.7831289 0.7821250 0.7811829
## [22] 0.7803680 0.7796745 0.7790821 0.7785507 0.7780645 0.7776579 0.7773076
## [29] 0.7769670 0.7766061 0.7762183 0.7758683 0.7755752 0.7753335 0.7751331
## [36] 0.7749673 0.7748375 0.7747401 0.7746698 0.7746088 0.7745580 0.7745132
## [43] 0.7744731 0.7744365 0.7744090 0.7743851 0.7743627 0.7743417 0.7743137
## [50] 0.7742775 0.7742361 0.7741572 0.7740726 0.7739959 0.7739332 0.7738777
## [57] 0.7738309 0.7737931 0.7737645 0.7737415 0.7737202 0.7736960 0.7736730
## [64] 0.7736527 0.7736363 0.7736223 0.7736116 0.7736033 0.7735973 0.7735931
## [71] 0.7735910
```

```
#So, to get the 'loss' value at lambda == lambda.1se
glm1s_cv$cvm [ which(glm1s_cv$lambda == glm1s_cv$lambda.1se) ]
```

```
## [1] 0.7773076
```

```
#PREDICTIONS on Train
```

```
glmPred1s_1=predict(glm1s_cv,data.matrix(xDTrn), s="lambda.min" )

glmPred1s_1p=predict(glm1s_cv,data.matrix(xDTrn), s="lambda.min", type="response" ) #gives the probabilities
# gives the the ln(p/(1-p)) values
#i.e. the values of w1*x1 + ...+w2*x2

# AUC using default type.measure = "deviance"
predsauc <- prediction(glmPred1s_1p, lcdfTrn$loan_status, label.ordering = c("Charged Off", "Fully Paid"))
aucPerf <- performance(predsauc, "auc")
aucPerf@y.values
```

```
## [[1]]
## [1] 0.6929972
```

```
#PREDICTIONS on Test
```

```
glmPredls_1_Tst=predict ( glm1s_cv,data.matrix(xDTst), s="lambda.min" )
```

```
glmPredls_1p_Tst=predict(glm1s_cv,data.matrix(xDTst), s="lambda.min", type="response" ) #gives the prob values
```

```
# gives the the  $\ln(p/(1-p))$  values
```

```
#i.e. the values of  $w_1*x_1 + \dots + w_2*x_2$ 
```

```
# AUC using default type.measure = "deviance"
```

```
predsauc_Tst <- prediction(glmPredls_1p_Tst, lcdfTst$loan_status, label.ordering = c("Charged Off", "Fully Paid"))
```

```
aucPerf_Tst <- performance(predsauc, "auc")
```

```
aucPerf_Tst@y.values
```

```
## [[1]]
```

```
## [1] 0.6929972
```

Experiment with Ridge regularization as a variable selection method.

```
#Use Ridge regularization ( $\alpha = 0$ )
```

```
glm1s_cv_L2<- cv.glmnet(data.matrix(xDTrn), yTrn, family="binomial", alpha = 0)
```


```
glm1s_cv_L2$lambda.min
```

```
## [1] 0.007635409
```

```
glm1s_cv_L2$lambda.1se
```

```
## [1] 0.2386612
```

```
plot(glm1s_cv_L2)
```



```
coef(glmnet_cv_L2, s = glmnet_cv_L2$lambda.min)
```

```
## 42 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)    3.483800e+00
## loan_amnt      -1.094290e-06
## int_rate       -4.191654e-02
## installment    -1.338318e-04
## grade          -1.070678e-01
## sub_grade      -2.802382e-02
## emp_length      1.141786e-02
## home_ownership -6.030660e-02
## annual_inc     -1.944640e-07
## purpose         3.599765e-03
## dti            -1.556590e-02
## initial_list_status -2.379598e-02
## total_rev_hi_lim  2.220878e-06
## acc_open_past_24mths -4.325442e-02
## avg_cur_bal     1.795154e-06
## bc_open_to_buy   1.253395e-06
## bc_util         -2.557541e-03
## chargeoff_within_12_mths -4.185435e-02
## delinq_amnt      2.594939e-05
## mo_sin_old_il_acct -9.960001e-05
## mo_sin_old_rev_tl_op  2.502301e-04
## mo_sin_rcnt_rev_tl_op -2.315311e-05
## mo_sin_rcnt_tl     2.741097e-03
## mort_acc         2.678016e-02
## mths_since_recent_bc -1.698771e-04
## mths_since_recent_inq  3.610931e-03
## num_bc_sats      -1.624826e-02
## num_bc_tl        -1.991454e-02
## num_il_tl         3.147844e-03
## num_op_rev_tl     -1.416594e-02
## num_rev_accts     9.194141e-03
## num_sats          4.837194e-03
## num_tl_30dpd      -2.030509e-01
## pct_tl_nvr_dlq     1.417863e-03
## tax_liens         -3.212956e-02
## tot_hi_cred_lim    4.972057e-07
## total_bal_ex_mort  -1.451971e-06
## total_bc_limit     4.584778e-06
## total_il_high_credit_limit 1.673504e-06
## propSatisBankcardAccts -5.312934e-02
## prop_OpAccts_to_TotAccts  1.936422e-01
## propLoanAmt_to_AnnInc -5.692460e-01
```

```
coef(glm1s_cv_L2, s = glm1s_cv_L2$lambda.1se)
```

```
## 42 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  2.873218e+00
## loan_amnt    -3.652773e-07
## int_rate     -2.736051e-02
## installment -4.858538e-05
## grade        -8.660296e-02
## sub_grade    -1.810468e-02
## emp_length    5.049587e-03
## home_ownership -3.509874e-02
## annual_inc    3.728905e-07
## purpose      -3.244608e-03
## dti           -6.565991e-03
## initial_list_status -2.553814e-02
## total_rev_hi_lim  8.137283e-07
## acc_open_past_24mths -1.847017e-02
## avg_cur_bal    1.893942e-06
## bc_open_to_buy  2.178659e-06
## bc_util        -1.250212e-03
## chargeoff_within_12_mths -9.273238e-03
## delinq_amnt    5.700583e-06
## mo_sin_old_il_acct -1.410239e-05
## mo_sin_old_rev_tl_op  2.018724e-04
## mo_sin_rcnt_rev_tl_op  8.876395e-04
## mo_sin_rcnt_tl    2.432139e-03
## mort_acc        1.493114e-02
## mths_since_recent_bc -3.200099e-05
## mths_since_recent_inq  1.956409e-03
## num_bc_sats     -4.399570e-03
## num_bc_tl       -2.487589e-03
## num_il_tl       -1.114332e-04
## num_op_rev_tl   -3.617182e-03
## num_rev_accts   -4.871645e-04
## num_sats        -8.986079e-04
## num_tl_30dpd    -6.352455e-02
## pct_tl_nvr_dlq   3.508673e-04
## tax_liens       -1.127259e-02
## tot_hi_cred_lim  2.027922e-07
## total_bal_ex_mort  8.237045e-08
## total_bc_limit   1.486659e-06
## total_il_high_credit_limit 1.557624e-07
## propSatisBankcardAccts -1.407273e-02
## prop_OpAccts_to_TotAccts -4.296392e-02
## propLoanAmt_to_AnnInc -3.709809e-01
```

```
#find the index of the best lambda
which(glm1s_cv_L2$lambda == glm1s_cv_L2$lambda.1se)
```

```
## [1] 63
```

```
glm1s_cv_L2$glmnet.fit$dev.ratio[which(glm1s_cv_L2$lambda == glm1s_cv_L2$lambda.1se) ]
```




```
## [1] 0.05971116
```

```
glm1s_cv_L2$glmnet.fit
```


```
##
## Call:  glmnet(x = data.matrix(xDTrn), y = yTrn, family = "binomial",      alpha = 0)
##
##      Df %Dev Lambda
## 1    41 0.00 76.350
## 2    41 0.10 69.570
## 3    41 0.11 63.390
## 4    41 0.12 57.760
## 5    41 0.13 52.630
## 6    41 0.14 47.950
## 7    41 0.15 43.690
## 8    41 0.17 39.810
## 9    41 0.18 36.270
## 10   41 0.20 33.050
## 11   41 0.22 30.120
## 12   41 0.24 27.440
## 13   41 0.27 25.000
## 14   41 0.29 22.780
## 15   41 0.32 20.760
## 16   41 0.35 18.910
## 17   41 0.38 17.230
## 18   41 0.41 15.700
## 19   41 0.45 14.310
## 20   41 0.49 13.040
## 21   41 0.54 11.880
## 22   41 0.59 10.820
## 23   41 0.64  9.862
## 24   41 0.70  8.985
## 25   41 0.76  8.187
## 26   41 0.83  7.460
## 27   41 0.90  6.797
## 28   41 0.97  6.193
## 29   41 1.06  5.643
## 30   41 1.14  5.142
## 31   41 1.24  4.685
## 32   41 1.34  4.269
## 33   41 1.45  3.890
## 34   41 1.56  3.544
## 35   41 1.68  3.229
## 36   41 1.81  2.942
## 37   41 1.94  2.681
## 38   41 2.08  2.443
## 39   41 2.22  2.226
## 40   41 2.37  2.028
## 41   41 2.53  1.848
## 42   41 2.69  1.684
## 43   41 2.86  1.534
## 44   41 3.03  1.398
## 45   41 3.21  1.274
## 46   41 3.38  1.161
## 47   41 3.56  1.057
## 48   41 3.74  0.964
## 49   41 3.92  0.878
## 50   41 4.10  0.800
```

```
## 51 41 4.27 0.729
## 52 41 4.45 0.664
## 53 41 4.62 0.605
## 54 41 4.78 0.551
## 55 41 4.94 0.502
## 56 41 5.10 0.458
## 57 41 5.24 0.417
## 58 41 5.38 0.380
## 59 41 5.52 0.346
## 60 41 5.64 0.316
## 61 41 5.76 0.288
## 62 41 5.87 0.262
## 63 41 5.97 0.239
## 64 41 6.07 0.218
## 65 41 6.15 0.198
## 66 41 6.23 0.180
## 67 41 6.31 0.164
## 68 41 6.37 0.150
## 69 41 6.44 0.137
## 70 41 6.49 0.124
## 71 41 6.54 0.113
## 72 41 6.59 0.103
## 73 41 6.63 0.094
## 74 41 6.67 0.086
## 75 41 6.70 0.078
## 76 41 6.73 0.071
## 77 41 6.76 0.065
## 78 41 6.78 0.059
## 79 41 6.80 0.054
## 80 41 6.82 0.049
## 81 41 6.84 0.045
## 82 41 6.86 0.041
## 83 41 6.87 0.037
## 84 41 6.89 0.034
## 85 41 6.90 0.031
## 86 41 6.91 0.028
## 87 41 6.92 0.026
## 88 41 6.93 0.023
## 89 41 6.94 0.021
## 90 41 6.95 0.019
## 91 41 6.95 0.018
## 92 41 6.96 0.016
## 93 41 6.97 0.015
## 94 41 6.97 0.013
## 95 41 6.98 0.012
## 96 41 6.99 0.011
## 97 41 6.99 0.010
## 98 41 7.00 0.009
## 99 41 7.00 0.008
## 100 41 7.01 0.008
```

```
plot(glm_12$glmnet.fit)
```



```
plot(glm1s_cv_L2$glmnet.fit, xvar="lambda")
```



```
#as.matrix(coef(glm_ls_cv_L2, s = glm_ls_cv_L2$lambda.min))
#as.matrix(coef(glm_ls_cv_L2, s = glm_ls_cv_L2$lambda.1se))
```

```
#the lambda values used are in glm_ls_cv_L2$lambda
glm_ls_cv_L2$lambda
```

```
## [1] 76.354092223 69.571005116 63.390508773 57.759070690 52.627914045
## [6] 47.952595214 43.692618822 39.811086991 36.274379750 33.051863832
## [11] 30.115627347 27.440238018 25.002522903 22.781367680 20.757533765
## [16] 18.913491677 17.233269206 15.702313068 14.307362854 13.036336172
## [21] 11.878223997 10.822995316 9.861510242 8.985440852 8.187199053
## [26] 7.459870854 6.797156489 6.193315840 5.643118731 5.141799617
## [31] 4.685016310 4.268812374 3.889582849 3.544043029 3.229200014
## [36] 2.942326785 2.680938583 2.442771389 2.225762311 2.028031722
## [41] 1.847866974 1.683707565 1.534131625 1.397843600 1.273663027
## [46] 1.160514315 1.057417421 0.963479372 0.877886520 0.799897501
## [51] 0.728836812 0.664088958 0.605093126 0.551338320 0.502358943
## [56] 0.457730760 0.417067220 0.380016118 0.346256534 0.315496057
## [61] 0.287468257 0.261930370 0.238661198 0.217459195 0.198140719
## [66] 0.180538444 0.164499906 0.149886188 0.136570712 0.124438146
## [71] 0.113383404 0.103310735 0.094132894 0.085770387 0.078150783
## [76] 0.071208082 0.064882153 0.059118201 0.053866303 0.049080969
## [81] 0.044720751 0.040747883 0.037127953 0.033829608 0.030824279
## [86] 0.028085935 0.025590858 0.023317437 0.021245980 0.019358546
## [91] 0.017638786 0.016071805 0.014644030 0.013343095 0.012157731
## [96] 0.011077672 0.010093562 0.009196878 0.008379852 0.007635409
```

```
# and the cross-validation 'loss' at each lambda is in glmLs_cv_L2$cvm
glmLs_cv_L2$cvm
```

```
## [1] 0.8307081 0.8299185 0.8298323 0.8297472 0.8296541 0.8295522 0.8294408
## [8] 0.8293189 0.8291856 0.8290399 0.8288808 0.8287070 0.8285174 0.8283105
## [15] 0.8280850 0.8278392 0.8275717 0.8272806 0.8269641 0.8266203 0.8262473
## [22] 0.8258431 0.8254141 0.8249423 0.8244328 0.8238834 0.8232920 0.8226564
## [29] 0.8219745 0.8212445 0.8204645 0.8196330 0.8187487 0.8178107 0.8168184
## [36] 0.8157717 0.8146711 0.8135176 0.8123128 0.8110590 0.8097592 0.8084172
## [43] 0.8070374 0.8056249 0.8041856 0.8027259 0.8012527 0.7997733 0.7982954
## [50] 0.7968267 0.7953750 0.7939478 0.7925523 0.7911954 0.7898833 0.7886214
## [57] 0.7874145 0.7862663 0.7851797 0.7841569 0.7831988 0.7823057 0.7814773
## [64] 0.7807121 0.7800076 0.7793631 0.7787752 0.7782408 0.7777566 0.7773193
## [71] 0.7769255 0.7765716 0.7762544 0.7759706 0.7757172 0.7754913 0.7752901
## [78] 0.7751111 0.7749520 0.7748107 0.7746852 0.7745736 0.7744744 0.7743862
## [85] 0.7743075 0.7742371 0.7741743 0.7741177 0.7740662 0.7740198 0.7739775
## [92] 0.7739387 0.7739029 0.7738693 0.7738377 0.7738086 0.7737808 0.7737544
## [99] 0.7737292 0.7737056
```

```
#So, to get the 'loss' value at lambda == lambda.1se
glmLs_cv_L2$cvm [ which(glmLs_cv_L2$lambda == glmLs_cv_L2$lambda.1se) ]
```

```
## [1] 0.7814773
```

```
#PREDICTIONS on Trn
```

```
glmPredls_1_L2=predict ( glm1s_cv_L2,data.matrix(xDTrn), s="lambda.min" )
```

```
glmPredls_1p_L2=predict(glm1s_cv_L2,data.matrix(xDTrn), s="lambda.min", type="response" ) #gives  
the prob values
```

```
# gives the the  $\ln(p/(1-p))$  values
```

```
#i.e. the values of  $w1*x1 + \dots + w2*x2$ 
```

```
# AUC using default type.measure = "deviance"
```

```
preds_L2 <- prediction(glmPredls_1p_L2, lcdfTrn$loan_status, label.ordering = c("Charged Off", "Fully Paid"))
```

```
aucPerf_L2 <- performance(preds_L2, "auc")
```

```
aucPerf_L2@y.values
```

```
## [[1]]
```

```
## [1] 0.6922717
```

```
#PREDICTIONS on Tst
```

```
glmPredls_1_L2_Tst=predict(glm1s_cv_L2,data.matrix(xDTst), s="lambda.min" )
```

```
glmPredls_1p_L2_Tst=predict(glm1s_cv_L2,data.matrix(xDTst), s="lambda.min", type="response" ) #gi  
ves the prob values
```

```
# gives the the  $\ln(p/(1-p))$  values
```

```
#i.e. the values of  $w1*x1 + \dots + w2*x2$ 
```

```
# AUC using default type.measure = "deviance"
```

```
preds_L2_Tst <- prediction(glmPredls_1p_L2_Tst, lcdfTst$loan_status, label.ordering = c("Charged  
Off", "Fully Paid"))
```

```
aucPerf_L2_Tst <- performance(preds_L2_Tst, "auc")
```

```
aucPerf_L2_Tst@y.values
```

```
## [[1]]
```

```
## [1] 0.6967046
```

```
#glm models with different Alpha values
```

```
# alpha default, 1
```

```
glmRet_cv<- cv.glmnet(data.matrix(xDTrn), lcdfTrn$actualReturn, family="gaussian")
```

```
plot(glmRet_cv)
```

```
glmRet_cv$lambda.min
```

```
## [1] 5.194606e-06
```

```
glmRet_cv$lambda.1se
```

```
## [1] 0.002646204
```

```
coef(glmRet_cv, s = glmRet_cv$lambda.min) %>% tidy()
```

```
## Warning: 'tidy.dgCMatrix' is deprecated.  
## See help("Deprecated")
```

```
## Warning: 'tidy.dgTMatrix' is deprecated.  
## See help("Deprecated")
```



##		row	column	value
## 1	(Intercept)	1		2.551692e-02
## 2	loan_amnt	1		7.711433e-08
## 3	int_rate	1		6.733179e-03
## 4	installment	1		-3.558124e-06
## 5	grade	1		-4.877115e-03
## 6	sub_grade	1		-1.558048e-03
## 7	emp_length	1		4.944113e-04
## 8	home_ownership	1		-2.042390e-03
## 9	annual_inc	1		-2.768122e-08
## 10	purpose	1		-4.498478e-05
## 11	dti	1		-6.011703e-04
## 12	initial_list_status	1		3.474949e-04
## 13	total_rev_hi_lim	1		7.894354e-08
## 14	acc_open_past_24mths	1		-8.177517e-04
## 15	avg_cur_bal	1		4.116041e-08
## 16	bc_open_to_buy	1		-2.021713e-07
## 17	bc_util	1		-9.341336e-05
## 18	chargeoff_within_12_mths	1		-2.829240e-03
## 19	delinq_amnt	1		6.748106e-07
## 20	mo_sin_old_il_acct	1		-3.732481e-06
## 21	mo_sin_old_rev_tl_op	1		-1.037648e-06
## 22	mo_sin_rcnt_rev_tl_op	1		-2.339319e-05
## 23	mo_sin_rcnt_tl	1		6.906570e-06
## 24	mort_acc	1		8.831819e-04
## 25	mths_since_recent_bc	1		-8.402183e-06
## 26	mths_since_recent_inq	1		3.532926e-05
## 27	num_bc_sats	1		5.099128e-04
## 28	num_bc_tl	1		-1.035082e-03
## 29	num_il_tl	1		2.026320e-04
## 30	num_op_rev_tl	1		-7.447672e-04
## 31	num_rev_accts	1		6.797010e-04
## 32	num_sats	1		-1.631254e-04
## 33	num_tl_30dpd	1		-7.533966e-03
## 34	pct_tl_nvr_dlq	1		1.312548e-05
## 35	tax_liens	1		-1.649921e-03
## 36	tot_hi_cred_lim	1		4.796931e-09
## 37	total_bal_ex_mort	1		-8.295016e-08
## 38	total_bc_limit	1		1.215017e-07
## 39	total_il_high_credit_limit	1		1.116969e-07
## 40	propSatisBankcardAccts	1		-6.587370e-03
## 41	prop_OpAccts_to_TotAccts	1		9.566931e-03
## 42	propLoanAmt_to_AnnInc	1		-2.997153e-02


```
coef(glmRet_cv, s = glmRet_cv$lambda.1se) %>% tidy()
```

```
## Warning: 'tidy.dgCMatrix' is deprecated.
## See help("Deprecated")
```

```
## Warning: 'tidy.dgTMatrix' is deprecated.
## See help("Deprecated")
```

##		row	column	value
## 1	(Intercept)	1	3.950458e-02	
## 2	int_rate	1	1.404950e-03	
## 3	home_ownership	1	-7.944462e-05	
## 4	dti	1	-1.727085e-04	
## 5	avg_cur_bal	1	2.702681e-08	
## 6	mort_acc	1	1.862549e-04	
## 7	propLoanAmt_to_AnnInc	1	-9.786915e-03	

```
# alpha = 0
glmRet_cv_a0<- cv.glmnet(data.matrix(xDTrn), lcdfTrn$actualReturn, family="gaussian", alpha=0)
plot(glmRet_cv)
```



```
glmRet_cv_a0$lambda.min
```

```
## [1] 0.000736322
```

```
glmRet_cv_a0$lambda.1se
```

```
## [1] 0.4116636
```

```
coef(glmRet_cv_a0, s = glmRet_cv_a0$lambda.min) %>% tidy()
```

```
## Warning: 'tidy.dgCMatrix' is deprecated.
## See help("Deprecated")

## Warning: 'tidy.dgTMatrix' is deprecated.
## See help("Deprecated")
```

```
##           row column      value
## 1      (Intercept)      1 4.172457e-02
## 2      loan_amnt      1 8.442883e-08
## 3      int_rate      1 3.624728e-03
## 4      installment      1 -3.480015e-06
## 5      grade      1 -4.165564e-03
## 6      sub_grade      1 3.110412e-04
## 7      emp_length      1 4.930394e-04
## 8      home_ownership      1 -2.020220e-03
## 9      annual_inc      1 -2.623814e-08
## 10     purpose      1 -3.575762e-05
## 11     dti      1 -5.808739e-04
## 12     initial_list_status      1 3.045555e-04
## 13     total_rev_hi_lim      1 6.343752e-08
## 14     acc_open_past_24mths      1 -8.040540e-04
## 15     avg_cur_bal      1 3.612892e-08
## 16     bc_open_to_buy      1 -1.698464e-07
## 17     bc_util      1 -8.963615e-05
## 18     chargeoff_within_12_mths      1 -2.776782e-03
## 19     delinq_amnt      1 6.744389e-07
## 20     mo_sin_old_il_acct      1 -3.630714e-06
## 21     mo_sin_old_rev_tl_op      1 -1.121481e-06
## 22     mo_sin_rcnt_rev_tl_op      1 -2.550418e-05
## 23     mo_sin_rcnt_tl      1 8.269895e-06
## 24     mort_acc      1 8.729466e-04
## 25     mths_since_recent_bc      1 -8.219525e-06
## 26     mths_since_recent_inq      1 3.368706e-05
## 27     num_bc_sats      1 3.188200e-04
## 28     num_bc_tl      1 -8.790961e-04
## 29     num_il_tl      1 1.821432e-04
## 30     num_op_rev_tl      1 -6.159420e-04
## 31     num_rev_accts      1 5.808111e-04
## 32     num_sats      1 -1.470538e-04
## 33     num_tl_30dpd      1 -7.359863e-03
## 34     pct_tl_nvr_dlq      1 1.052302e-05
## 35     tax_liens      1 -1.685867e-03
## 36     tot_hi_cred_lim      1 5.221320e-09
## 37     total_bal_ex_mort      1 -6.406478e-08
## 38     total_bc_limit      1 1.043928e-07
## 39     total_il_high_credit_limit      1 9.207458e-08
## 40     propSatisBankcardAccts      1 -5.551993e-03
## 41     propOpAccts_to_TotAccts      1 8.017315e-03
## 42     propLoanAmt_to_AnnInc      1 -3.044067e-02
```

```
coef(glmRet_cv_a0, s = glmRet_cv_a0$lambda.1se) %>% tidy()
```

```
## Warning: 'tidy.dgCMatrix' is deprecated.
```


```
## See help("Deprecated")
```

```
## Warning: 'tidy.dgTMatrix' is deprecated.
```

```
## See help("Deprecated")
```

```
##           row column      value
## 1      (Intercept)      1 5.209616e-02
## 2      loan_amnt      1 -2.033871e-08
## 3      int_rate      1 2.568672e-04
## 4      installment      1 -3.136385e-07
## 5      grade      1 7.079644e-04
## 6      sub_grade      1 1.616860e-04
## 7      emp_length      1 9.138598e-05
## 8      home_ownership      1 -4.227668e-04
## 9      annual_inc      1 3.118059e-09
## 10     purpose      1 2.063172e-04
## 11     dti      1 -6.922019e-05
## 12     initial_list_status      1 2.435018e-04
## 13     total_rev_hi_lim      1 -2.012277e-09
## 14     acc_open_past_24mths      1 -1.163193e-05
## 15     avg_cur_bal      1 2.412000e-08
## 16     bc_open_to_buy      1 -1.124526e-08
## 17     bc_util      1 -4.001027e-06
## 18     chargeoff_within_12_mths      1 -8.015508e-05
## 19     delinq_amnt      1 1.552436e-07
## 20     mo_sin_old_il_acct      1 -8.516995e-08
## 21     mo_sin_old_rev_tl_op      1 -7.351034e-07
## 22     mo_sin_rcnt_rev_tl_op      1 -6.686794e-06
## 23     mo_sin_rcnt_tl      1 -1.276974e-05
## 24     mort_acc      1 2.110347e-04
## 25     mths_since_recent_bc      1 -2.158574e-07
## 26     mths_since_recent_inq      1 -8.759551e-06
## 27     num_bc_sats      1 -1.174914e-04
## 28     num_bc_tl      1 -5.357285e-05
## 29     num_il_tl      1 2.285276e-06
## 30     num_op_rev_tl      1 -6.221624e-05
## 31     num_rev_accts      1 -1.230585e-05
## 32     num_sats      1 -3.814338e-05
## 33     num_tl_30dpd      1 -2.255909e-04
## 34     pct_tl_nvr_dlq      1 -2.476893e-05
## 35     tax_liens      1 -1.782960e-04
## 36     tot_hi_cred_lim      1 1.876635e-09
## 37     total_bal_ex_mort      1 6.393171e-10
## 38     total_bc_limit      1 -8.612982e-09
## 39     total_il_high_credit_limit      1 1.055736e-09
## 40     propSatisBankcardAccts      1 -3.630662e-04
## 41     propOpAccts_to_TotAccts      1 -8.039163e-04
## 42     propLoanAmt_to_AnnInc      1 -4.832487e-03
```

```
# alpha = 0.2
glmRet_cv_a2<- cv.glmnet(data.matrix(xDTrn), lcdfTrn$actualReturn, family="gaussian", alpha=0.2)
plot(glmRet_cv_a2)
```



```
glmRet_cv_a2$lambda.min
```

```
## [1] 3.68161e-06
```

```
glmRet_cv_a2$lambda.1se
```

```
## [1] 0.01749063
```

```
coef(glmRet_cv_a2, s = glmRet_cv_a2$lambda.min) %>% tidy()
```

```
## Warning: 'tidy.dgCMatrix' is deprecated.
## See help("Deprecated")
```

```
## Warning: 'tidy.dgTMatrix' is deprecated.
## See help("Deprecated")
```

##		row	column	value
## 1	(Intercept)	1		2.376534e-02
## 2	loan_amnt	1		3.142668e-07
## 3	int_rate	1		7.001775e-03
## 4	installment	1		-1.059921e-05
## 5	grade	1		-4.696537e-03
## 6	sub_grade	1		-1.741281e-03
## 7	emp_length	1		4.961373e-04
## 8	home_ownership	1		-2.052012e-03
## 9	annual_inc	1		-2.813165e-08
## 10	purpose	1		-4.855901e-05
## 11	dti	1		-6.037904e-04
## 12	initial_list_status	1		3.597045e-04
## 13	total_rev_hi_lim	1		8.158254e-08
## 14	acc_open_past_24mths	1		-8.180106e-04
## 15	avg_cur_bal	1		4.200113e-08
## 16	bc_open_to_buy	1		-2.086637e-07
## 17	bc_util	1		-9.427909e-05
## 18	chargeoff_within_12_mths	1		-2.862435e-03
## 19	delinq_amnt	1		6.816663e-07
## 20	mo_sin_old_il_acct	1		-3.788225e-06
## 21	mo_sin_old_rev_tl_op	1		-1.090545e-06
## 22	mo_sin_rcnt_rev_tl_op	1		-2.396128e-05
## 23	mo_sin_rcnt_tl	1		8.344806e-06
## 24	mort_acc	1		8.887078e-04
## 25	mths_since_recent_bc	1		-8.548070e-06
## 26	mths_since_recent_inq	1		3.570044e-05
## 27	num_bc_sats	1		5.699533e-04
## 28	num_bc_tl	1		-1.074260e-03
## 29	num_il_tl	1		2.121088e-04
## 30	num_op_rev_tl	1		-7.685054e-04
## 31	num_rev_accts	1		7.023511e-04
## 32	num_sats	1		-1.783152e-04
## 33	num_tl_30dpd	1		-7.614860e-03
## 34	pct_tl_nvr_dlq	1		1.346322e-05
## 35	tax_liens	1		-1.652604e-03
## 36	tot_hi_cred_lim	1		4.702166e-09
## 37	total_bal_ex_mort	1		-8.589938e-08
## 38	total_bc_limit	1		1.229475e-07
## 39	total_il_high_credit_limit	1		1.146485e-07
## 40	propSatisBankcardAccts	1		-6.969269e-03
## 41	prop_OpAccts_to_TotAccts	1		1.011881e-02
## 42	propLoanAmt_to_AnnInc	1		-3.004989e-02


```
coef(glmRet_cv_a2, s = glmRet_cv_a2$lambda.1se) %>% tidy()
```

```
## Warning: 'tidy.dgCMatrix' is deprecated.  
## See help("Deprecated")
```

```
## Warning: 'tidy.dgTMatrix' is deprecated.  
## See help("Deprecated")
```

##		row	column	value
## 1	(Intercept)	1		4.327400e-02
## 2	int_rate	1		6.371532e-04
## 3	sub_grade	1		2.094240e-04
## 4	dti	1		-5.444355e-05
## 5	propLoanAmt_to_AnnInc	1		-2.340863e-03

```
# alpha = 0.5
glmRet_cv_a5<- cv.glmnet(data.matrix(xDTrn), lcdfTrn$actualReturn, family="gaussian", alpha=0.5)
plot(glmRet_cv_a5)
```



```
glmRet_cv_a5$lambda.min
```

```
## [1] 1.815544e-05
```

```
glmRet_cv_a5$lambda.1se
```

```
## [1] 0.005808411
```

```
coef(glmRet_cv_a5, s = glmRet_cv_a5$lambda.min) %>% tidy()
```

```
## Warning: 'tidy.dgCMatrix' is deprecated.
## See help("Deprecated")
```

```
## Warning: 'tidy.dgTMatrix' is deprecated.
## See help("Deprecated")
```

```
##           row column      value
## 1      (Intercept)      1 2.772876e-02
## 2      loan_amnt      1 3.054229e-08
## 3      int_rate      1 6.331441e-03
## 4      installment      1 -2.117222e-06
## 5      grade      1 -4.906894e-03
## 6      sub_grade      1 -1.296690e-03
## 7      emp_length      1 4.934160e-04
## 8      home_ownership      1 -2.036826e-03
## 9      annual_inc      1 -2.733392e-08
## 10     purpose      1 -4.208124e-05
## 11     dti      1 -5.981707e-04
## 12     initial_list_status      1 3.332394e-04
## 13     total_rev_hi_lim      1 7.613381e-08
## 14     acc_open_past_24mths      1 -8.160141e-04
## 15     avg_cur_bal      1 4.011380e-08
## 16     bc_open_to_buy      1 -1.951264e-07
## 17     bc_util      1 -9.260309e-05
## 18     chargeoff_within_12_mths      1 -2.790830e-03
## 19     delinq_amnt      1 6.685484e-07
## 20     mo_sin_old_il_acct      1 -3.674688e-06
## 21     mo_sin_old_rev_tl_op      1 -9.921886e-07
## 22     mo_sin_rcnt_rev_tl_op      1 -2.308049e-05
## 23     mo_sin_rcnt_tl      1 5.897328e-06
## 24     mort_acc      1 8.761491e-04
## 25     mths_since_recent_bc      1 -8.300803e-06
## 26     mths_since_recent_inq      1 3.492583e-05
## 27     num_bc_sats      1 4.625985e-04
## 28     num_bc_tl      1 -1.002729e-03
## 29     num_il_tl      1 1.930228e-04
## 30     num_op_rev_tl      1 -7.280963e-04
## 31     num_rev_accts      1 6.601328e-04
## 32     num_sats      1 -1.464486e-04
## 33     num_tl_30dpd      1 -7.445361e-03
## 34     pct_tl_nvr_dlq      1 1.272934e-05
## 35     tax_liens      1 -1.648265e-03
## 36     tot_hi_cred_lim      1 4.858040e-09
## 37     total_bal_ex_mort      1 -7.962659e-08
## 38     total_bc_limit      1 1.184235e-07
## 39     total_il_high_credit_limit      1 1.082874e-07
## 40     propSatisBankcardAccts      1 -6.259324e-03
## 41     prop_OpAccts_to_TotAccts      1 9.033403e-03
## 42     propLoanAmt_to_AnnInc      1 -3.003045e-02
```

```
coef(glmRet_cv_a5, s = glmRet_cv_a5$lambda.1se) %>% tidy()
```



```
## Warning: 'tidy.dgCMatrix' is deprecated.  
## See help("Deprecated")  
  
## Warning: 'tidy.dgTMatrix' is deprecated.  
## See help("Deprecated")
```

```
##           row column      value  
## 1      (Intercept)      1 4.022190e-02  
## 2         int_rate      1 1.249527e-03  
## 3            dti      1 -1.343694e-04  
## 4      avg_cur_bal      1 1.840565e-08  
## 5         mort_acc      1 8.777688e-05  
## 6 propLoanAmt_to_AnnInc      1 -7.451321e-03
```

```
sum(yTrn == 0)
```

```
## [1] 8268
```

```
sum(yTrn == 1)
```

```
## [1] 48447
```


```
1-sum(yTrn == 0)/length(yTrn)
```

```
## [1] 0.8542185
```

```
1-sum(yTrn == 1)/length(yTrn)
```

```
## [1] 0.1457815
```

```
wts <- if_else(yTrn == 0, 1-sum(yTrn == 0)/length(yTrn), 1-sum(yTrn == 1)/length(yTrn))  
  
# Lasso regularization, balanced weights, type.measure='deviance'  
glm1sw_cv<- cv.glmnet(data.matrix(xDTrn), yTrn, family= "binomial", weights = wts, alpha = 1)  
plot(glm1sw_cv)
```



```
#Prediction on Trn
glmPredls_balp=predict(glm1sw_cv,data.matrix(xDTrn), s="lambda.min", type="response" )

preds_auc_bal <- prediction(glmPredls_balp, lcdfTrn$loan_status, label.ordering = c("Charged Off"
, "Fully Paid"))
aucPerf_bal <- performance(preds_auc_bal, "auc")
aucPerf_bal@y.values
```

```
## [[1]]
## [1] 0.6930635
```

```
#Prediction on Tst
glmPredls_balp_Tst=predict(glm1sw_cv,data.matrix(xDTst), s="lambda.min", type="response" )

preds_auc_bal_Tst <- prediction(glmPredls_balp_Tst, lcdfTst$loan_status, label.ordering = c("Char
ged Off", "Fully Paid"))
aucPerf_bal_Tst <- performance(preds_auc_bal_Tst, "auc")
aucPerf_bal_Tst@y.values
```

```
## [[1]]
## [1] 0.6965404
```

We tried building a non-regularized model, but the resulting AUC is no better than the best regularized glmnet model we have so far. The AUC on Test data is 0.6886417.

```
#build glm model without regularization, using coefficient from the cross-validated model.  
glm1s_nzv_2 <- glm(yTrn ~ data.matrix(xDTrn %>% select(nzCoefVars)), family=binomial())
```

```
## Note: Using an external vector in selections is ambiguous.  
## i Use `all_of(nzCoefVars)` instead of `nzCoefVars` to silence this message.  
## i See <https://tidyselect.r-lib.org/reference/faq-external-vector.html>.  
## This message is displayed once per session.
```

```
summary(glm1s_nzv_2)
```

```
##
## Call:
## glm(formula = yTrn ~ data.matrix(xDTrn %>% select(nzCoefVars)),
##     family = binomial())
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7950   0.3430   0.4671   0.6019   1.5698
##
## Coefficients:
##                                     Estimate
## (Intercept)                        3.730e+00
## data.matrix(xDTrn %>% select(nzCoefVars))int_rate -4.446e-02
## data.matrix(xDTrn %>% select(nzCoefVars))grade -7.467e-02
## data.matrix(xDTrn %>% select(nzCoefVars))sub_grade -3.563e-02
## data.matrix(xDTrn %>% select(nzCoefVars))home_ownership -6.166e-02
## data.matrix(xDTrn %>% select(nzCoefVars))dti -1.291e-02
## data.matrix(xDTrn %>% select(nzCoefVars))acc_open_past_24mths -5.214e-02
## data.matrix(xDTrn %>% select(nzCoefVars))bc_open_to_buy 4.238e-06
## data.matrix(xDTrn %>% select(nzCoefVars))bc_util -2.457e-03
## data.matrix(xDTrn %>% select(nzCoefVars))mort_acc 2.380e-02
## data.matrix(xDTrn %>% select(nzCoefVars))mths_since_recent_inq 4.450e-03
## data.matrix(xDTrn %>% select(nzCoefVars))tot_hi_cred_lim 8.551e-07
## data.matrix(xDTrn %>% select(nzCoefVars))propLoanAmt_to_AnnInc -8.644e-01
##                                     Std. Error
## (Intercept)                        2.055e-01
## data.matrix(xDTrn %>% select(nzCoefVars))int_rate 3.476e-02
## data.matrix(xDTrn %>% select(nzCoefVars))grade 4.322e-02
## data.matrix(xDTrn %>% select(nzCoefVars))sub_grade 2.468e-02
## data.matrix(xDTrn %>% select(nzCoefVars))home_ownership 1.616e-02
## data.matrix(xDTrn %>% select(nzCoefVars))dti 1.487e-03
## data.matrix(xDTrn %>% select(nzCoefVars))acc_open_past_24mths 4.152e-03
## data.matrix(xDTrn %>% select(nzCoefVars))bc_open_to_buy 1.505e-06
## data.matrix(xDTrn %>% select(nzCoefVars))bc_util 5.839e-04
## data.matrix(xDTrn %>% select(nzCoefVars))mort_acc 8.056e-03
## data.matrix(xDTrn %>% select(nzCoefVars))mths_since_recent_inq 1.016e-03
## data.matrix(xDTrn %>% select(nzCoefVars))tot_hi_cred_lim 1.211e-07
## data.matrix(xDTrn %>% select(nzCoefVars))propLoanAmt_to_AnnInc 1.181e-01
##                                     z value Pr(>|z|)
## (Intercept)                        18.149 < 2e-16
## data.matrix(xDTrn %>% select(nzCoefVars))int_rate -1.279 0.200912
## data.matrix(xDTrn %>% select(nzCoefVars))grade -1.728 0.084066
## data.matrix(xDTrn %>% select(nzCoefVars))sub_grade -1.443 0.148921
## data.matrix(xDTrn %>% select(nzCoefVars))home_ownership -3.815 0.000136
## data.matrix(xDTrn %>% select(nzCoefVars))dti -8.682 < 2e-16
## data.matrix(xDTrn %>% select(nzCoefVars))acc_open_past_24mths -12.555 < 2e-16
## data.matrix(xDTrn %>% select(nzCoefVars))bc_open_to_buy 2.817 0.004854
## data.matrix(xDTrn %>% select(nzCoefVars))bc_util -4.208 2.57e-05
## data.matrix(xDTrn %>% select(nzCoefVars))mort_acc 2.954 0.003135
## data.matrix(xDTrn %>% select(nzCoefVars))mths_since_recent_inq 4.379 1.19e-05
## data.matrix(xDTrn %>% select(nzCoefVars))tot_hi_cred_lim 7.059 1.68e-12
## data.matrix(xDTrn %>% select(nzCoefVars))propLoanAmt_to_AnnInc -7.320 2.49e-13
##
## (Intercept) ***
```

```
## data.matrix(xDTrn %>% select(nzCoefVars))int_rate
## data.matrix(xDTrn %>% select(nzCoefVars))grade .
## data.matrix(xDTrn %>% select(nzCoefVars))sub_grade
## data.matrix(xDTrn %>% select(nzCoefVars))home_ownership ***
## data.matrix(xDTrn %>% select(nzCoefVars))dti ***
## data.matrix(xDTrn %>% select(nzCoefVars))acc_open_past_24mths ***
## data.matrix(xDTrn %>% select(nzCoefVars))bc_open_to_buy **
## data.matrix(xDTrn %>% select(nzCoefVars))bc_util ***
## data.matrix(xDTrn %>% select(nzCoefVars))mort_acc **
## data.matrix(xDTrn %>% select(nzCoefVars))mths_since_recent_inq ***
## data.matrix(xDTrn %>% select(nzCoefVars))tot_hi_cred_lim ***
## data.matrix(xDTrn %>% select(nzCoefVars))propLoanAmt_to_AnnInc ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 47110 on 56714 degrees of freedom
## Residual deviance: 43930 on 56702 degrees of freedom
## AIC: 43956
##
## Number of Fisher Scoring iterations: 5
```

```
tidy(glm1s_nzv_2)
```

```
## # A tibble: 13 x 5
##   term                                estimate std.error statistic  p.value
##   <chr>                                <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)                        3.73e+0    2.06e-1    18.1  1.31e-73
## 2 data.matrix(xDTrn %>% select(nzCoef~ -4.45e-2    3.48e-2    -1.28  2.01e- 1
## 3 data.matrix(xDTrn %>% select(nzCoef~ -7.47e-2    4.32e-2    -1.73  8.41e- 2
## 4 data.matrix(xDTrn %>% select(nzCoef~ -3.56e-2    2.47e-2    -1.44  1.49e- 1
## 5 data.matrix(xDTrn %>% select(nzCoef~ -6.17e-2    1.62e-2    -3.82  1.36e- 4
## 6 data.matrix(xDTrn %>% select(nzCoef~ -1.29e-2    1.49e-3    -8.68  3.90e-18
## 7 data.matrix(xDTrn %>% select(nzCoef~ -5.21e-2    4.15e-3   -12.6  3.71e-36
## 8 data.matrix(xDTrn %>% select(nzCoef~  4.24e-6    1.50e-6     2.82  4.85e- 3
## 9 data.matrix(xDTrn %>% select(nzCoef~ -2.46e-3    5.84e-4    -4.21  2.57e- 5
## 10 data.matrix(xDTrn %>% select(nzCoef~  2.38e-2    8.06e-3     2.95  3.14e- 3
## 11 data.matrix(xDTrn %>% select(nzCoef~  4.45e-3    1.02e-3     4.38  1.19e- 5
## 12 data.matrix(xDTrn %>% select(nzCoef~  8.55e-7    1.21e-7     7.06  1.68e-12
## 13 data.matrix(xDTrn %>% select(nzCoef~ -8.64e-1    1.18e-1    -7.32  2.49e-13
```

```
#PREDICTIONS on Trn
```

```
glmPredls_1_nonreg= predict(glm1s_nzv_2,xDTrn)
```

```
glmPredls_1p_nonreg=predict(glm1s_nzv_2, xDTrn, type="response" ) #gives the prob values
```

```
# AUC using non regularized glm
```

```
preds_nonreg <- prediction(glmPredls_1p_nonreg, lcdfTrn$loan_status, label.ordering = c("Charged Off", "Fully Paid"))
```

```
aucPerf_nonreg <- performance(preds_nonreg, "auc")
```

```
aucPerf_nonreg@y.values
```

```
## [[1]]
## [1] 0.6891976
```

```
#PREDICTIONS on Tst
glmPredls_1_nonreg_Tst= predict(glm1s_nzv_2,xDTst)
```

```
## Warning: 'newdata' had 24307 rows but variables found have 56715 rows
```

```
glmPredls_1p_nonreg_Tst=predict(glm1s_nzv_2, xDTst, type="response" ) #gives the prob values
```

```
## Warning: 'newdata' had 24307 rows but variables found have 56715 rows
```

```
# AUC using non regularized glm
preds_nonreg_Tst <- prediction(glmPredls_1p_nonreg_Tst, lcdfTrn$loan_status, label.ordering = c(
"Charged Off", "Fully Paid"))
aucPerf_nonreg_Tst <- performance(preds_nonreg, "auc")
aucPerf_nonreg_Tst@y.values
```

```
## [[1]]
## [1] 0.6891976
```

```
#Use no cross-validation model
glm1s_1 <- glmnet(data.matrix(xDTrn), yTrn, family="binomial", lambda = glm1s_cv$lambda.1se )
glm1s_1
```

```
##
## Call:  glmnet(x = data.matrix(xDTrn), y = yTrn, family = "binomial",      lambda = glm1s_cv$lambda.1se)
##
##      Df %Dev   Lambda
## 1 12 6.48 0.006193
```

## 1B. Experiment with loss and link function

Link function is used to fit our target variable to the requirement for the scale of glm. A glm, like any regression model, fits the target variable to a scale from negative infinity to positive infinity. In this case, a prediction of two classes (either “Fully Paid” or “Charged Off”) does not fulfil this criteria. Therefore, the family parameter helps to set the condition of our model to fit the glm algorithm.

For the binomial family, the valid link functions are logit, probit, cauchit. We set the family to binomial and we use the link function logit. Logit is taking the log of the odds of class 1 for the dependent variable loan status.

The loss function we choose will be tuned in the parameter “type.measure”. Since this is a logistic regression, there are three possible loss functions we can use.


We experimented on the loss functions: deviance, auc, and class. Please refer to the first table for the results for which one gave the best AUC performance. Type.measure = deviance constantly gives the best AUC performance despite other parameters. Type.measure = auc gives similar results, but deviance still gives the best results on test data.

For models using type.measure = deviance (default), please refer to the answers for question 1A.

Link Function source: <https://www.r-bloggers.com/2018/10/generalized-linear-models-understanding-the-link-function/> (<https://www.r-bloggers.com/2018/10/generalized-linear-models-understanding-the-link-function/>)

```
#experiment with type.measure (LOSS FUNCTION), with alpha = 1

#type.measure = "auc"
glm1s_cv_auc<- cv.glmnet(data.matrix(xDTrn), yTrn, family="binomial", type.measure = "auc")
plot(glm1s_cv_auc)
```



```
#PREDICTIONS with AUC model with Train Data
glmPredls_1_auc=predict(glmmls_cv_auc,data.matrix(xDTrn), s="lambda.min" )

glmPredls_1p_auc=predict(glmmls_cv_auc,data.matrix(xDTrn), s="lambda.min", type="response" ) #give
s the prob values
# gives the the  $\ln(p/(1-p))$  values
#i.e. the values of  $w1*x1 + \dots + w2*x2$ 

preds_auc <- prediction(glmPredls_1p_auc, lcdfTrn$loan_status, label.ordering = c("Charged Off",
"Fully Paid"))
aucPerf_auc <- performance(preds_auc, "auc")
aucPerf_auc@y.values
```

```
## [[1]]
## [1] 0.6929972
```

```
#PREDICTIONS with AUC model on Test Data
glmPredls_1_auc_Tst=predict(glmmls_cv_auc,data.matrix(xDTst), s="lambda.min" )

glmPredls_1p_auc_Tst=predict(glmmls_cv_auc,data.matrix(xDTst), s="lambda.min", type="response" ) #
gives the prob values
# gives the the  $\ln(p/(1-p))$  values
#i.e. the values of  $w1*x1 + \dots + w2*x2$ 

preds_auc_Tst <- prediction(glmPredls_1p_auc_Tst, lcdfTst$loan_status, label.ordering = c("Charge
d Off", "Fully Paid"))
aucPerf_auc_Tst <- performance(preds_auc_Tst, "auc")
aucPerf_auc_Tst@y.values
```

```
## [[1]]
## [1] 0.6966821
```

```
#the labmda values used
glmmls_cv_auc$lambda
```



```
## [1] 0.0763540922 0.0695710051 0.0633905088 0.0577590707 0.0526279140
## [6] 0.0479525952 0.0436926188 0.0398110870 0.0362743798 0.0330518638
## [11] 0.0301156273 0.0274402380 0.0250025229 0.0227813677 0.0207575338
## [16] 0.0189134917 0.0172332692 0.0157023131 0.0143073629 0.0130363362
## [21] 0.0118782240 0.0108229953 0.0098615102 0.0089854409 0.0081871991
## [26] 0.0074598709 0.0067971565 0.0061933158 0.0056431187 0.0051417996
## [31] 0.0046850163 0.0042688124 0.0038895828 0.0035440430 0.0032292000
## [36] 0.0029423268 0.0026809386 0.0024427714 0.0022257623 0.0020280317
## [41] 0.0018478670 0.0016837076 0.0015341316 0.0013978436 0.0012736630
## [46] 0.0011605143 0.0010574174 0.0009634794 0.0008778865 0.0007998975
## [51] 0.0007288368 0.0006640890 0.0006050931 0.0005513383 0.0005023589
## [56] 0.0004577308 0.0004170672 0.0003800161 0.0003462565 0.0003154961
## [61] 0.0002874683 0.0002619304 0.0002386612 0.0002174592 0.0001981407
## [66] 0.0001805384 0.0001644999 0.0001498862 0.0001365707 0.0001244381
## [71] 0.0001133834
```


```
# and the cross-validation 'loss' at each lambda
glm1s_cv_auc$cvm
```

```
## [1] 0.5652220 0.6732776 0.6732776 0.6732776 0.6732479 0.6732479 0.6732212
## [8] 0.6732212 0.6733284 0.6734491 0.6734535 0.6734432 0.6734499 0.6739759
## [15] 0.6751585 0.6764871 0.6784778 0.6800665 0.6812982 0.6823313 0.6833596
## [22] 0.6842020 0.6849100 0.6854603 0.6859668 0.6864247 0.6867836 0.6871123
## [29] 0.6874334 0.6877421 0.6880999 0.6884384 0.6887115 0.6889153 0.6890645
## [36] 0.6891757 0.6892723 0.6893261 0.6893524 0.6893606 0.6893761 0.6893781
## [43] 0.6893595 0.6893628 0.6893622 0.6893838 0.6894101 0.6894356 0.6894767
## [50] 0.6895196 0.6896009 0.6897022 0.6897960 0.6899101 0.6900071 0.6900846
## [57] 0.6901569 0.6902041 0.6902456 0.6902868 0.6903189 0.6903685 0.6904137
## [64] 0.6904563 0.6904842 0.6905095 0.6905321 0.6905518 0.6905686 0.6905758
## [71] 0.6905896
```

```
#So, to get the 'loss' value at lambda == lambda.1se
glm1s_cv_auc$cvm [ which(glm1s_cv_auc$lambda == glm1s_cv_auc$lambda.1se) ]
```

```
## [1] 0.6877421
```

```
#type.measure = "class"
glm1s_cv_class<- cv.glmnet(data.matrix(xDTrn), yTrn, family="binomial", type.measure = "class")
plot(glm1s_cv_class)
```



```
#PREDICTIONS with "class" model on Train
glmPredls_1_class=predict(glmLs_cv_class,data.matrix(xDTrn), s="lambda.min" )

glmPredls_1p_class=predict(glmLs_cv_class,data.matrix(xDTrn), s="lambda.min", type="response" ) #
gives the prob values
# gives the the  $\ln(p/(1-p))$  values
#i.e. the values of  $w1*x1 + \dots + w2*x2$ 

preds_class <- prediction(glmPredls_1p_class, lcdfTrn$loan_status, label.ordering = c("Charged Of
f", "Fully Paid"))
aucPerf_class <- performance(preds_class, "auc")
aucPerf_class@y.values
```

```
## [[1]]
## [1] 0.6855892
```

```
#PREDICTIONS with "class" model on Test
glmPredls_1_class_Tst=predict(glmmls_cv_class,data.matrix(xDTst), s="lambda.min" )

glmPredls_1p_class_Tst=predict(glmmls_cv_class,data.matrix(xDTst), s="lambda.min", type="response"
) #gives the prob values
# gives the the ln(p/(1-p)) values
#i.e. the values of w1*x1 + ...+w2*x2

preds_class_Tst <- prediction(glmPredls_1p_class_Tst, lcdfTst$loan_status, label.ordering = c("Charged Off", "Fully Paid"))
aucPerf_class_Tst <- performance(preds_class_Tst, "auc")
aucPerf_class_Tst@y.values
```

```
## [[1]]
## [1] 0.6900923
```

```
#the labmda values used
glmmls_cv_class$lambda
```

```
## [1] 0.0763540922 0.0695710051 0.0633905088 0.0577590707 0.0526279140
## [6] 0.0479525952 0.0436926188 0.0398110870 0.0362743798 0.0330518638
## [11] 0.0301156273 0.0274402380 0.0250025229 0.0227813677 0.0207575338
## [16] 0.0189134917 0.0172332692 0.0157023131 0.0143073629 0.0130363362
## [21] 0.0118782240 0.0108229953 0.0098615102 0.0089854409 0.0081871991
## [26] 0.0074598709 0.0067971565 0.0061933158 0.0056431187 0.0051417996
## [31] 0.0046850163 0.0042688124 0.0038895828 0.0035440430 0.0032292000
## [36] 0.0029423268 0.0026809386 0.0024427714 0.0022257623 0.0020280317
## [41] 0.0018478670 0.0016837076 0.0015341316 0.0013978436 0.0012736630
## [46] 0.0011605143 0.0010574174 0.0009634794 0.0008778865 0.0007998975
## [51] 0.0007288368 0.0006640890 0.0006050931 0.0005513383 0.0005023589
## [56] 0.0004577308 0.0004170672 0.0003800161 0.0003462565 0.0003154961
## [61] 0.0002874683 0.0002619304 0.0002386612 0.0002174592 0.0001981407
## [66] 0.0001805384 0.0001644999 0.0001498862 0.0001365707 0.0001244381
## [71] 0.0001133834
```

```
# and the cross-validation 'loss' at each lambda
glmmls_cv_class$cvm
```


```
## [1] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [8] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [15] 0.1457992 0.1457815 0.1457815 0.1458168 0.1457463 0.1457463 0.1456934
## [22] 0.1456934 0.1456757 0.1457286 0.1457463 0.1457992 0.1457992 0.1458697
## [29] 0.1458168 0.1458168 0.1458344 0.1458873 0.1458873 0.1459050 0.1459402
## [36] 0.1459402 0.1459579 0.1459402 0.1460460 0.1460460 0.1460284 0.1460108
## [43] 0.1460460 0.1459755 0.1459579 0.1459579 0.1459931 0.1459755 0.1459931
## [50] 0.1460108 0.1459755 0.1459755 0.1459579 0.1459931 0.1459931 0.1459755
## [57] 0.1459755 0.1459050 0.1459226 0.1459402 0.1459579 0.1459755 0.1459755
## [64] 0.1459931 0.1460108 0.1460108 0.1460108 0.1460108 0.1460108 0.1460460
## [71] 0.1460460
```

```
#So, to get the 'loss' value at lambda == lambda.1se
glm1s_cv_class$cvm [ which(glm1s_cv_class$lambda == glm1s_cv_class$lambda.1se) ]
```

```
## [1] 0.1457815
```

```
#experiment with type.measure, with alpha = 0
```

```
#type.measure = "auc"
glm1s_cv_auc_L2<- cv.glmnet(data.matrix(xDTrn), yTrn, family="binomial", type.measure = "auc", al
pha = 0)
plot(glm1s_cv_auc_L2)
```



```
#PREDICTIONS with AUC model on Trn
glmPredls_1_auc_L2=predict(glmLs_cv_auc_L2,data.matrix(xDTrn), s="lambda.min" )

glmPredls_1p_auc_L2=predict(glmLs_cv_auc_L2,data.matrix(xDTrn), s="lambda.min", type="response" )
#gives the prob values
# gives the the ln(p/(1-p)) values
#i.e. the values of w1*x1 + ...+w2*x2

preds_auc_L2 <- prediction(glmPredls_1p_auc_L2, lcdfTrn$loan_status, label.ordering = c("Charged
Off", "Fully Paid"))
aucPerf_auc_L2 <- performance(preds_auc_L2, "auc")
aucPerf_auc_L2@y.values
```

```
## [[1]]
## [1] 0.6922717
```

```
#PREDICTIONS with AUC model on Test
glmPredls_1_auc_L2_Tst=predict(glmLs_cv_auc_L2,data.matrix(xDTst), s="lambda.min" )

glmPredls_1p_auc_L2_Tst=predict(glmLs_cv_auc_L2,data.matrix(xDTst), s="lambda.min", type="response" ) #gives the prob values
# gives the the ln(p/(1-p)) values
#i.e. the values of w1*x1 + ...+w2*x2

preds_auc_L2_Tst <- prediction(glmPredls_1p_auc_L2_Tst, lcdfTst$loan_status, label.ordering = c(
"Charged Off", "Fully Paid"))
aucPerf_auc_L2_Tst <- performance(preds_auc_L2_Tst, "auc")
aucPerf_auc_L2_Tst@y.values
```

```
## [[1]]
## [1] 0.6967046
```

```
#the lambda values used
glmLs_cv_auc_L2$lambda
```

```
## [1] 76.354092223 69.571005116 63.390508773 57.759070690 52.627914045
## [6] 47.952595214 43.692618822 39.811086991 36.274379750 33.051863832
## [11] 30.115627347 27.440238018 25.002522903 22.781367680 20.757533765
## [16] 18.913491677 17.233269206 15.702313068 14.307362854 13.036336172
## [21] 11.878223997 10.822995316 9.861510242 8.985440852 8.187199053
## [26] 7.459870854 6.797156489 6.193315840 5.643118731 5.141799617
## [31] 4.685016310 4.268812374 3.889582849 3.544043029 3.229200014
## [36] 2.942326785 2.680938583 2.442771389 2.225762311 2.028031722
## [41] 1.847866974 1.683707565 1.534131625 1.397843600 1.273663027
## [46] 1.160514315 1.057417421 0.963479372 0.877886520 0.799897501
## [51] 0.728836812 0.664088958 0.605093126 0.551338320 0.502358943
## [56] 0.457730760 0.417067220 0.380016118 0.346256534 0.315496057
## [61] 0.287468257 0.261930370 0.238661198 0.217459195 0.198140719
## [66] 0.180538444 0.164499906 0.149886188 0.136570712 0.124438146
## [71] 0.113383404 0.103310735 0.094132894 0.085770387 0.078150783
## [76] 0.071208082 0.064882153 0.059118201 0.053866303 0.049080969
## [81] 0.044720751 0.040747883 0.037127953 0.033829608 0.030824279
## [86] 0.028085935 0.025590858 0.023317437 0.021245980 0.019358546
## [91] 0.017638786 0.016071805 0.014644030 0.013343095 0.012157731
## [96] 0.011077672 0.010093562 0.009196878 0.008379852 0.007635409
```


```
# and the cross-validation 'loss' at each lambda
glm1s_cv_auc_L2$cvm
```

```
## [1] 0.6249439 0.6831316 0.6831393 0.6831525 0.6831627 0.6831756 0.6831876
## [8] 0.6832014 0.6832215 0.6832385 0.6832564 0.6832798 0.6833025 0.6833281
## [15] 0.6833535 0.6833861 0.6834195 0.6834565 0.6834924 0.6835349 0.6835774
## [22] 0.6836271 0.6836997 0.6837559 0.6838273 0.6838957 0.6839615 0.6840391
## [29] 0.6841252 0.6842132 0.6843104 0.6844130 0.6845126 0.6846247 0.6847394
## [36] 0.6848594 0.6849866 0.6851208 0.6852596 0.6853999 0.6855484 0.6857008
## [43] 0.6858475 0.6860078 0.6861622 0.6863107 0.6864538 0.6866084 0.6867627
## [50] 0.6869139 0.6870596 0.6872012 0.6873429 0.6874919 0.6876446 0.6877857
## [57] 0.6879213 0.6880399 0.6881624 0.6882725 0.6883757 0.6884911 0.6885996
## [64] 0.6887079 0.6888015 0.6888881 0.6889882 0.6890746 0.6891607 0.6892355
## [71] 0.6893152 0.6893935 0.6894561 0.6895224 0.6895997 0.6896573 0.6897233
## [78] 0.6897674 0.6898126 0.6898678 0.6899250 0.6899739 0.6900240 0.6900657
## [85] 0.6901069 0.6901516 0.6901890 0.6902198 0.6902578 0.6902932 0.6903361
## [92] 0.6903658 0.6903922 0.6904121 0.6904446 0.6904727 0.6905065 0.6905350
## [99] 0.6905595 0.6905860
```

```
#So, to get the 'loss' value at lambda == lambda.1se
glm1s_cv_auc_L2$cvm [ which(glm1s_cv_auc_L2$lambda == glm1s_cv_auc_L2$lambda.1se) ]
```

```
## [1] 0.6877857
```

```
#type.measure = "class"
glm1s_cv_class_L2<- cv.glmnet(data.matrix(xDTrn), yTrn, family="binomial", type.measure = "class"
, alpha = 0)
plot(glm1s_cv_class_L2)
```



```
#PREDICTIONS with "class" model on Trn
glmPredls_1_class_L2=predict(glmLs_cv_class_L2,data.matrix(xDTrn), s="lambda.min" )

glmPredls_1p_class_L2=predict(glmLs_cv_class_L2,data.matrix(xDTrn), s="lambda.min", type="response" ) #gives the prob values
# gives the the  $\ln(p/(1-p))$  values
#i.e. the values of  $w_1*x_1 + \dots + w_2*x_2$ 

preds_class_L2 <- prediction(glmPredls_1p_class_L2, lcdfTrn$loan_status, label.ordering = c("Charged Off", "Fully Paid"))
aucPerf_class_L2 <- performance(preds_class_L2, "auc")
aucPerf_class_L2@y.values
```

```
## [[1]]
## [1] 0.6903123
```

```
#PREDICTIONS with "class" model on Tst
glmPredls_1_class_L2_Tst=predict(glmLs_cv_class_L2,data.matrix(xDTst), s="lambda.min" )

glmPredls_1p_class_L2_Tst=predict(glmLs_cv_class_L2,data.matrix(xDTst), s="lambda.min", type="response" ) #gives the prob values
# gives the the ln(p/(1-p)) values
#i.e. the values of w1*x1 + ...+w2*x2

preds_class_L2_Tst <- prediction(glmPredls_1p_class_L2_Tst, lcdfTst$loan_status, label.ordering =
c("Charged Off", "Fully Paid"))
aucPerf_class_L2_Tst <- performance(preds_class_L2_Tst, "auc")
aucPerf_class_L2_Tst@y.values
```

```
## [[1]]
## [1] 0.6955733
```

```
#the labmda values used
glmLs_cv_class_L2$lambda
```

```
## [1] 76.354092223 69.571005116 63.390508773 57.759070690 52.627914045
## [6] 47.952595214 43.692618822 39.811086991 36.274379750 33.051863832
## [11] 30.115627347 27.440238018 25.002522903 22.781367680 20.757533765
## [16] 18.913491677 17.233269206 15.702313068 14.307362854 13.036336172
## [21] 11.878223997 10.822995316 9.861510242 8.985440852 8.187199053
## [26] 7.459870854 6.797156489 6.193315840 5.643118731 5.141799617
## [31] 4.685016310 4.268812374 3.889582849 3.544043029 3.229200014
## [36] 2.942326785 2.680938583 2.442771389 2.225762311 2.028031722
## [41] 1.847866974 1.683707565 1.534131625 1.397843600 1.273663027
## [46] 1.160514315 1.057417421 0.963479372 0.877886520 0.799897501
## [51] 0.728836812 0.664088958 0.605093126 0.551338320 0.502358943
## [56] 0.457730760 0.417067220 0.380016118 0.346256534 0.315496057
## [61] 0.287468257 0.261930370 0.238661198 0.217459195 0.198140719
## [66] 0.180538444 0.164499906 0.149886188 0.136570712 0.124438146
## [71] 0.113383404 0.103310735 0.094132894 0.085770387 0.078150783
## [76] 0.071208082 0.064882153 0.059118201 0.053866303 0.049080969
## [81] 0.044720751 0.040747883 0.037127953 0.033829608 0.030824279
## [86] 0.028085935 0.025590858 0.023317437 0.021245980 0.019358546
## [91] 0.017638786 0.016071805 0.014644030 0.013343095 0.012157731
## [96] 0.011077672 0.010093562 0.009196878 0.008379852 0.007635409
```

```
# and the cross-validation 'loss' at each lambda
glmLs_cv_class_L2$cvm
```



```
## [1] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [8] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [15] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [22] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [29] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [36] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [43] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [50] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [57] 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815 0.1457815
## [64] 0.1457815 0.1457815 0.1458168 0.1457815 0.1457639 0.1457110 0.1456757
## [71] 0.1456757 0.1456405 0.1456581 0.1456757 0.1457286 0.1457992 0.1458168
## [78] 0.1458521 0.1458344 0.1458344 0.1458344 0.1458873 0.1459050 0.1459226
## [85] 0.1459226 0.1459226 0.1460108 0.1459931 0.1460108 0.1459755 0.1460108
## [92] 0.1460460 0.1460460 0.1460460 0.1460637 0.1460989 0.1460637 0.1460637
## [99] 0.1460813 0.1460989
```

```
#So, to get the 'loss' value at lambda == lambda.1se
glm1s_cv_class_L2$cvm [ which(glm1s_cv_class_L2$lambda == glm1s_cv_class_L2$lambda.1se) ]
```

```
## [1] 0.1457815
```

## 1C

Compared to random forest and xgboost models that we created last time, the results are pretty similar. The AUC scores on test data are all around .69. Ranger is especially accurate on the training set, scoring around .79. However, the ranger's random forest score for test data is no different to the other models.

## 1D

The variable importance of the glm model conflicts with those of ranger and xgboost (even with rpart and C5.0 decision trees). The different models have different algorithms, hence the different variable importance. To some extent, multicollinearity can help to explain the differences. In linear models, correlated independent variables can turn out to have weaker coefficients than they are supposed to have independently (IDS 570, Linear Regression Lecture Material).

Installment is ranked as the 11th most important in ranger, and 9th in XGboost, but it is ranked 29th in glmnet. This variable is at the top of the multicollinearity list, highly correlated with loan\_amnt. This evidence alludes to our hypothesis that is multicollinearity exists, the coefficients do not reflect the true relationships of the independent variable with the target variable.

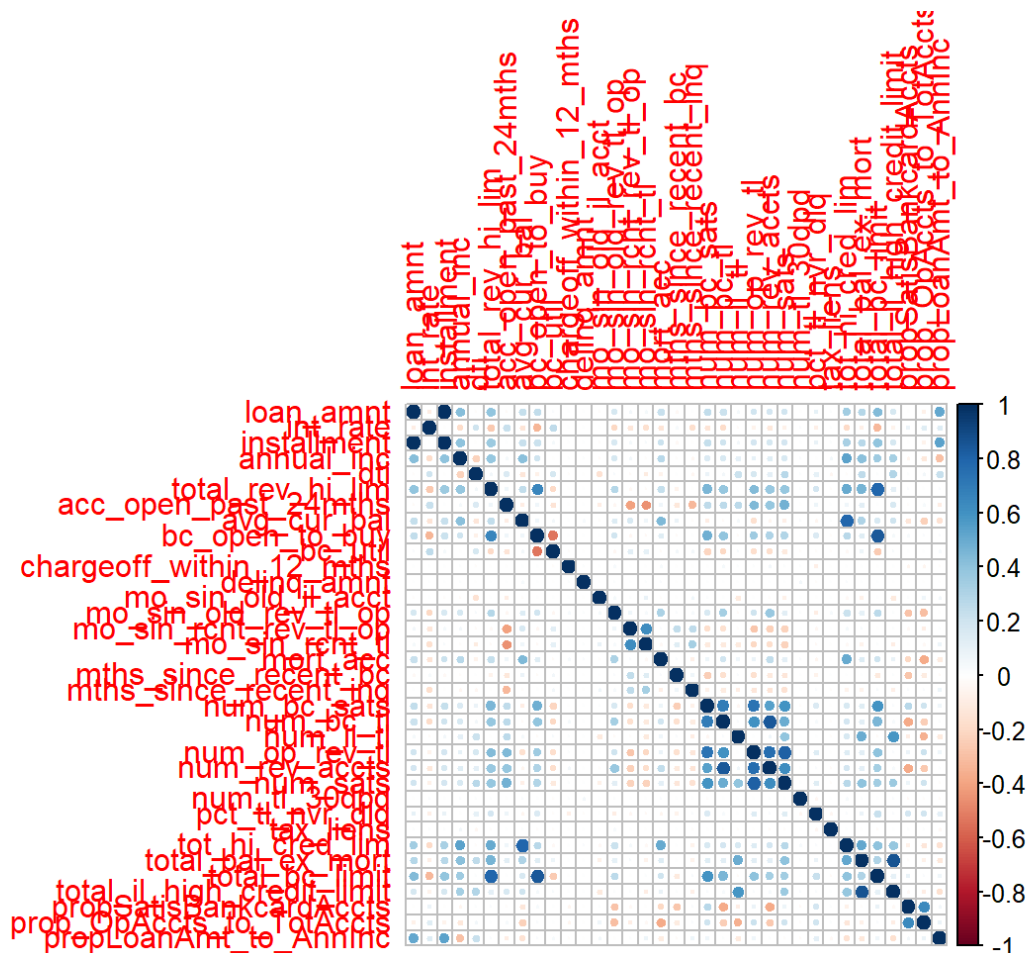
For the complete comparison of the variable importance between different models, please refer to our Appendix: Variable importance for different models.

Source: <https://www.r-bloggers.com/2020/07/comparing-variable-importance-functions-for-modeling/> (<https://www.r-bloggers.com/2020/07/comparing-variable-importance-functions-for-modeling/>)

```
library(corrplot)
```

```
## corrplot 0.84 loaded
```


```
xCorr <- xDTrn %>% select_if(is.numeric) %>% cor()
corrplot(xCorr, method="circle")
```



```
corrTH = 0.5
xCorr[upper.tri(xCorr, diag=TRUE)] <- NA #set the upper-diagonal values to NA
xCorr <- as.data.frame(as.table(xCorr))
xCorr <- na.omit(xCorr) #remove the rows corresponding to NA values
xCorr_th <- xCorr %>% filter(abs(Freq) > corrTH) #remove the rows with abs(values) < corrTH
xCorr_th <- xCorr_th[order(-abs(xCorr_th$Freq)),] #order by the corr values

#Convert back to matrix form to use with corrPlot
xCorrMat <- xCorr_th %>% pivot_wider(names_from = Var2, values_from = Freq)
xCorrMat <- column_to_rownames(xCorrMat, var="Var1") #convert first column to rownames

corrplot(as.matrix(xCorrMat), is.corr=FALSE, na.label=" ", method="circle")
```



## 1E

We tried to balance the two classes represented in the target variable when building the model. We are using 3 methods: - under sampling (US): under sample the majority class, in this case “Fully Paid”. The resulting training dataset has less data points (about 16,000 data points) - over sampling (OS): over sample the majority class, in this case “Charged Off”. The resulting training dataset has more data points (about 96,000 data points) - both (under sampling and over sampling) (BS): combining both under sampling and over sampling, resulting in moderate number of data points (56,000 data points)

We build models with identical parameters ( $\alpha = 1$ , type.measure = deviance) except for the training data set. The reason for this combination is that this combination resulted in the best AUC scores according to our experiment.

After running the code repeatedly, we have very similar results between OS, US, and BS. They only vary within less than .01 decimal points, and the order of AUC scores changes. This imply that the models we build are data dependent, and no method is ultimately better than the other.

We did not find evidence that concludes more data points caused by oversampling results in better models. This is intuitively true because copying/deleting existing examples do not enrich our dataset.

However, in the event of availability of new real data, more data is preferable. “Ideally, once you have more training examples you’ll have lower test-error (variance of the model decrease, meaning we are less overfitting), but theoretically, more data doesn’t always mean you will have more accurate model since high bias models will not benefit from more training examples.”

Source: <https://stats.stackexchange.com/questions/31249/what-impact-does-increasing-the-training-data-have-on-the-overall-system-accurac#> (<https://stats.stackexchange.com/questions/31249/what-impact-does-increasing-the-training-data-have-on-the-overall-system-accurac#>):~:text=Ideally%2C%20once%20you%20have%20more,benefit%20from%20more%20training%20examples.

```
library(ROSE)
```

```
## Warning: package 'ROSE' was built under R version 4.0.4
```

```
## Loaded ROSE 0.0-3
```

```
#Balancing the (training) data -- with over- and under-sampling
```

```
dim(lcdfTrn)
```

```
## [1] 56715    46
```

```
dim(lcdfTst)
```

```
## [1] 24307    46
```

```
us_lcdfTrn<-ovun.sample(loan_status~., data = as.data.frame(lcdfTrn), na.action = na.pass, method  
="under", p=0.5)$data  
dim(us_lcdfTrn)
```

```
## [1] 16672    46
```

```
#dim(LcdfTrn)  
us_lcdfTrn %>% group_by(loan_status) %>% tally()
```

```
## # A tibble: 2 x 2  
##   loan_status      n  
##   <fct>         <int>  
## 1 Fully Paid    8404  
## 2 Charged Off  8268
```

```
os_lcdfTrn<-ovun.sample(loan_status~., data = as.data.frame(lcdfTrn), na.action = na.pass, method  
="over", p=0.5)$data  
dim(os_lcdfTrn)
```

```
## [1] 96807    46
```

```
os_lcdfTrn %>% group_by(loan_status) %>% tally()
```

```
## # A tibble: 2 x 2
##   loan_status      n
##   <fct>         <int>
## 1 Fully Paid  48447
## 2 Charged Off 48360
```

```
bs_lcdfTrn<-ovun.sample(loan_status~., data = as.data.frame(lcdfTrn), na.action = na.pass, method
="both", p=0.5)$data
dim(bs_lcdfTrn)
```

```
## [1] 56715    46
```

```
bs_lcdfTrn %>% group_by(loan_status) %>% tally()
```

```
## # A tibble: 2 x 2
##   loan_status      n
##   <fct>         <int>
## 1 Fully Paid  28404
## 2 Charged Off 28311
```

```

yTrn_bs<-factor(if_else(bs_lcdfTrn$loan_status=="Fully Paid", '1', '0') )
yTrn_os<-factor(if_else(os_lcdfTrn$loan_status=="Fully Paid", '1', '0') )
yTrn_us<-factor(if_else(us_lcdfTrn$loan_status=="Fully Paid", '1', '0') )

bs_lcdfTrn_del<-bs_lcdfTrn %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_p
ymnt)
os_lcdfTrn_del<-os_lcdfTrn %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_p
ymnt)
us_lcdfTrn_del<-us_lcdfTrn %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_p
ymnt)

glm1s_cv_bs<- cv.glmnet(data.matrix(bs_lcdfTrn_del), yTrn_bs, family="binomial")
glm1s_cv_os<- cv.glmnet(data.matrix(os_lcdfTrn_del), yTrn_os, family="binomial")
glm1s_cv_us<- cv.glmnet(data.matrix(us_lcdfTrn_del), yTrn_us, family="binomial")

#BS
glmPred1s_1_bs=predict ( glm1s_cv_bs,data.matrix(bs_lcdfTrn_del), s="lambda.min" )

glmPred1s_1p_bs=predict(glm1s_cv_bs,data.matrix(bs_lcdfTrn_del), s="lambda.min", type="response"
)

predsauc_bs <- prediction(glmPred1s_1p_bs, bs_lcdfTrn$loan_status, label.ordering = c("Charged Of
f", "Fully Paid"))
aucPerf_bs <- performance(predsauc_bs, "auc")
aucPerf_bs@y.values

```

```

## [[1]]
## [1] 0.694004

```

```

#OS
glmPred1s_1_os=predict ( glm1s_cv_os,data.matrix(os_lcdfTrn_del), s="lambda.min" )

glmPred1s_1p_os=predict(glm1s_cv_os,data.matrix(os_lcdfTrn_del), s="lambda.min", type="response"
)

predsauc_os <- prediction(glmPred1s_1p_os, os_lcdfTrn$loan_status, label.ordering = c("Charged Of
f", "Fully Paid"))
aucPerf_os <- performance(predsauc_os, "auc")
aucPerf_os@y.values

```

```

## [[1]]
## [1] 0.6915689

```

```
#US
glmPredls_1_us=predict ( glm1s_cv_us,data.matrix(us_lcdfTrn_del), s="lambda.min" )

glmPredls_1p_us=predict(glm1s_cv_us,data.matrix(us_lcdfTrn_del), s="lambda.min", type="response"
)

predsauc_us <- prediction(glmPredls_1p_us, us_lcdfTrn$loan_status, label.ordering = c("Charged Of
f", "Fully Paid"))
aucPerf_us <- performance(predsauc_us, "auc")
aucPerf_us@y.values
```

```
## [[1]]
## [1] 0.6949481
```

## Q2: Create Models to Predict Actual Return

The variable “actualReturn” is the ratio of money made by the investors out of their total investment. It is created by subtracting the funded amount (amount invested by investors) from the total payment (amount paid to lending club by borrowers), divided by the funded amount. So far it only reflects the annual return, so we multiply it by the actual term of the loan (how long it takes for the loan to be paid back). This is true only if there was any attempt to pay back the loan at all (the actualTerm is > 0). If the loan was never paid back, this calculation does not reflect the actualReturn, so we assign the value 0 (there is no return/money made from that loan). The actualReturn variable DOES NOT account for Lending Club’s service fees (2-6% from the borrowers, and 1% from the investors).

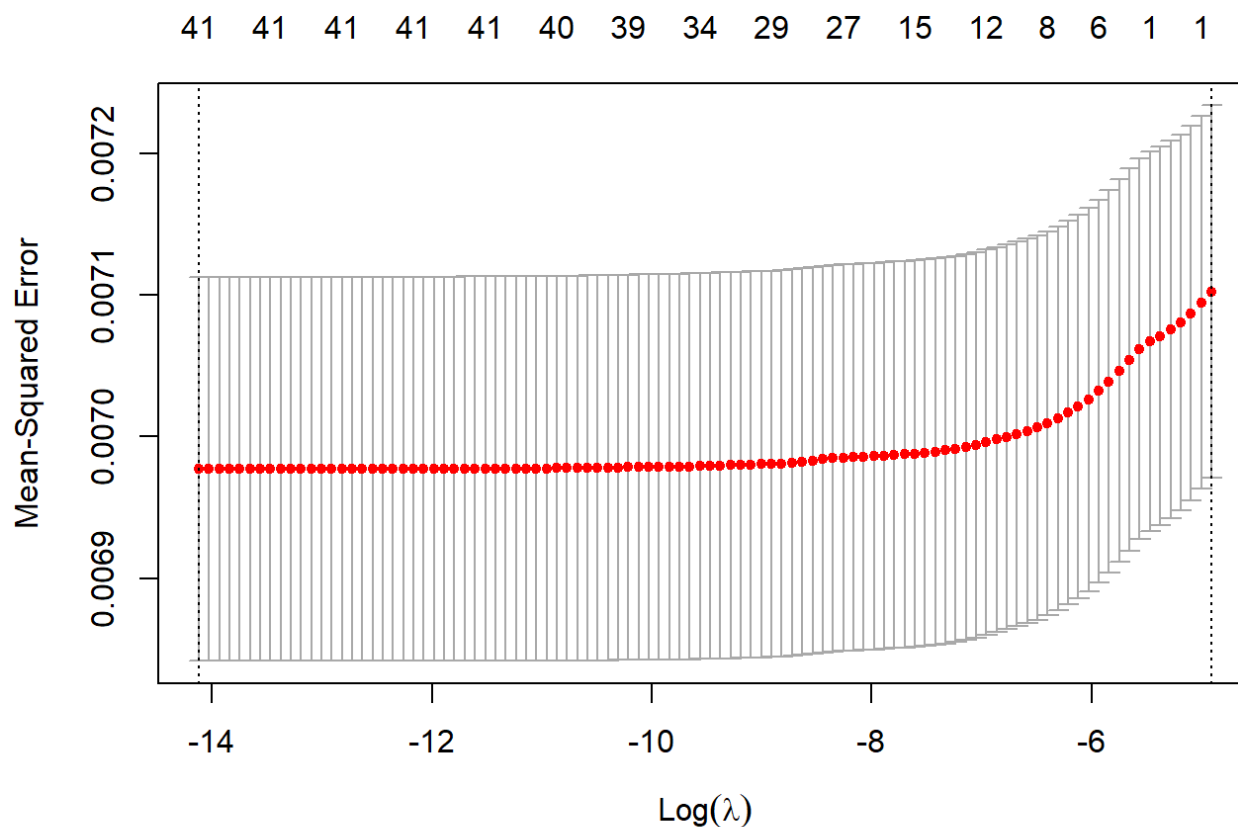
“Lending Club makes money by charging borrowers an origination fee and investors a service fee. The size of the origination fee depends on the credit grade and ranges to be 2%-6.0% of the loan amount. The size of the service fee is 1% on all amounts the borrower pays.” Source: - <https://www.investopedia.com/articles/investing/092315/7-best-peertopeer-lending-websites.asp> (<https://www.investopedia.com/articles/investing/092315/7-best-peertopeer-lending-websites.asp>) - <https://en.wikipedia.org/wiki/LendingClub> (<https://en.wikipedia.org/wiki/LendingClub>)

To calculate the true income for the investors, the actualReturn should be subtracted by 1% if the value is not 0 (at least 1 pay back was made on the loan).

```
library(glmnet)
library(tidyverse)
library(lubridate)

#Decile table with glmnet for Actual Return

xD<- lcdfTrn %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_pymnt)
glmRet_cv<- cv.glmnet(data.matrix(xD), lcdfTrn$actualReturn, family="gaussian")
plot(glmRet_cv)
```



```
#plot ( predict(glmRet_cv, data.matrix(lcdfTrn %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_pymnt)), s="lambda.min" ), lcdfTrn$actualReturn, main = "GLM Actual Return Model on Training Data")
```

```
#plot ( predict(glmRet_cv, data.matrix(lcdfTst %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_pymnt)), s="lambda.min" ), lcdfTst$actualReturn, main = "GLM Actual Return Model on Test Data")
```

```
#On Train Data
```

```
predRet_Trn_glm <- lcdfTrn %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
mutate(predRet= predict(glmRet_cv, data.matrix(lcdfTrn %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_pymnt)),s="lambda.min" ))
predRet_Trn_glm <- predRet_Trn_glm %>% mutate(tile=ntile(-predRet, 10))
predRet_Trn_glm %>% group_by(tile) %>% summarise(count=n(), avgpredRet=mean(predRet), numDefaults=sum(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(actualTerm), totA=sum(grade=="A"),
totB=sum(grade=="B" ), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F" ) )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```



```
## # A tibble: 10 x 14
##   tile count avgpredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>     <dbl>         <int>     <dbl>   <dbl>   <dbl>   <dbl> <int> <int>
## 1     1     1  5672     0.0731         1074    0.0727 -0.333   0.367   2.09    57   779
## 2     2     2  5672     0.0634          916    0.0658 -0.322   0.314   2.11   199  1409
## 3     3     3  5672     0.0590          924    0.0589 -0.333   0.307   2.19   398  1601
## 4     4     4  5672     0.0555          873    0.0563 -0.333   0.243   2.17   714  1743
## 5     5     5  5672     0.0525          850    0.0518 -0.311   0.259   2.21  1074  1809
## 6     6     6  5671     0.0496          799    0.0488 -0.333   0.225   2.24  1392  1928
## 7     7     7  5671     0.0467          766    0.0454 -0.323   0.248   2.27  1847  1834
## 8     8     8  5671     0.0435          720    0.0428 -0.312   0.225   2.28  2313  1842
## 9     9     9  5671     0.0396          699    0.0386 -0.323   0.210   2.31  2815  1720
## 10    10    10  5671     0.0324          647    0.0344 -0.323   0.229   2.39  3474  1664
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

*#on Test Data*

```
predRet_Tst_glm <- lcdfTst %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
mutate(predRet= predict(glmRet_cv, data.matrix(lcdfTst %>% select(-loan_status, -actualTerm, -ann
Ret, -actualReturn, -total_pymnt)),
s="lambda.min" ))
```

```
predRet_Tst_glm <- predRet_Tst_glm %>% mutate(tile=ntile(-predRet, 10))
```

```
predRet_Tst_glm %>% group_by(tile) %>% summarise(count=n(), avgpredRet=mean(predRet), numDefaults
=sum(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(act
ualTerm), totA=sum(grade=="A"),
totB=sum(grade=="B" ), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum
(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgpredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>     <dbl>         <int>     <dbl>   <dbl>   <dbl>   <dbl> <int> <int>
## 1     1     1  2431     0.0733          449    0.0733 -0.322   0.444   2.08    14   345
## 2     2     2  2431     0.0635          396    0.0667 -0.333   0.334   2.13    71   583
## 3     3     3  2431     0.0589          386    0.0601 -0.333   0.277   2.14   190   712
## 4     4     4  2431     0.0555          393    0.0538 -0.322   0.397   2.19   310   785
## 5     5     5  2431     0.0524          358    0.0516 -0.312   0.262   2.23   441   805
## 6     6     6  2431     0.0495          374    0.0462 -0.309   0.223   2.24   604   815
## 7     7     7  2431     0.0467          330    0.0465 -0.312   0.248   2.27   805   783
## 8     8     8  2430     0.0436          311    0.0416 -0.312   0.238   2.31   993   788
## 9     9     9  2430     0.0397          273    0.0408 -0.322   0.208   2.31  1204   749
## 10    10    10  2430     0.0322          289    0.0341 -0.323   0.230   2.37  1487   705
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

We created the decile table for Training and Test data using the glmnet model. Our model is able to distinguish the loans with the highest actualReturn, which are concentrated on the top deciles. As we go down the deciles, the average predicted return decreases.

However, this model does not avoid defaults very well. Even though the Average Actual Return for the top deciles are the highest, there are too many defaults. Combining the Actual Return and Loan Status model is necessary.

```
library(ranger)
```

```
rfModel_Ret <- ranger(actualReturn ~., data=subset(lcdfTrn, select=-c(loan_status, annRet, actualTerm, total_pymnt)), importance = "permutation", num.trees = 200)
```


```
## Growing trees.. Progress: 98%. Estimated remaining time: 0 seconds.  
## Computing permutation importance.. Progress: 52%. Estimated remaining time: 28 seconds.
```

```
rfPredRet_trn<- predict(rfModel_Ret, lcdfTrn)  
sqrt(mean( (rfPredRet_trn$predictions - lcdfTrn$actualReturn)^2))
```

```
## [1] 0.03759537
```

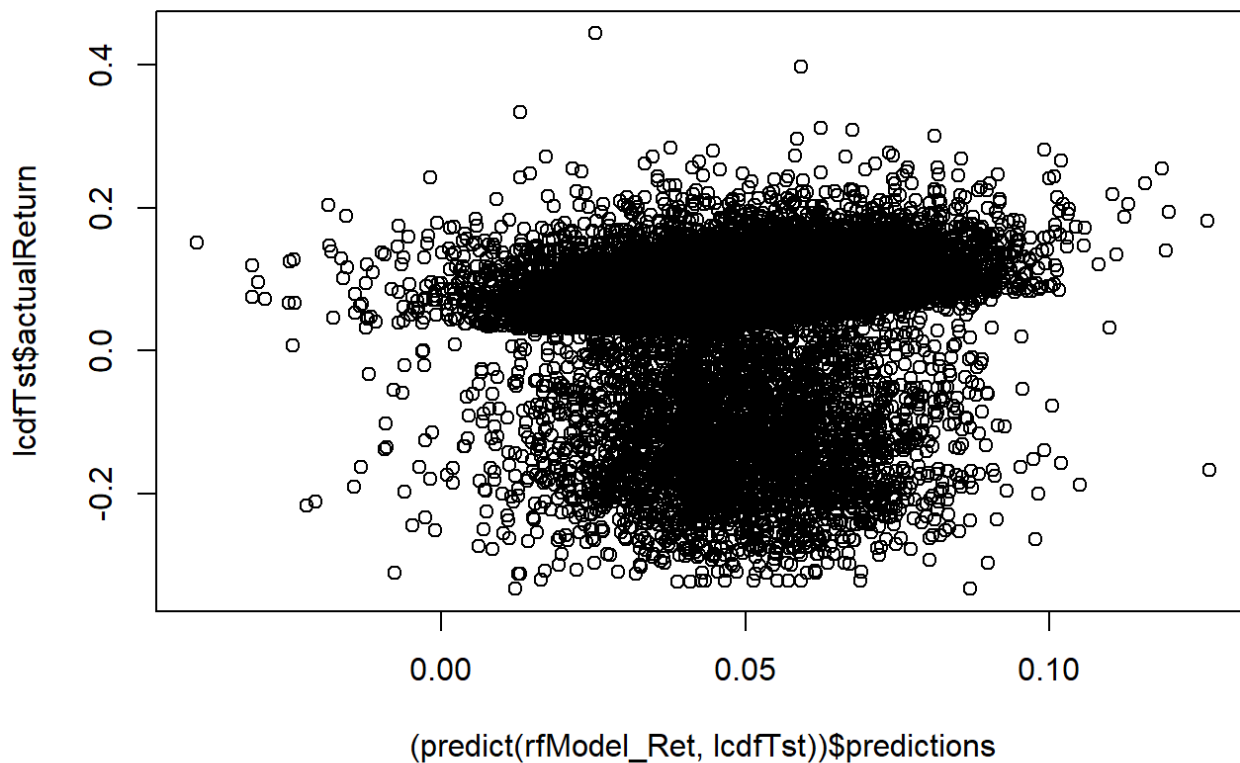
```
#sqrt(mean( (predict(rfModel_Ret, lcdfTst))$predictions - lcdfTst$actualReturn)^2))  
plot ( (predict(rfModel_Ret, lcdfTrn))$predictions, lcdfTrn$actualReturn, main = "RF Actual Return Model on Training Data")
```

### RF Actual Return Model on Training Data



```
plot ( (predict(rfModel_Ret, lcdfTst))$predictions, lcdfTst$actualReturn, main = "RF Actual Return Model on Test Data")
```

## RF Actual Return Model on Test Data



*#Performance by deciles on Trn*

```
predRet_Trn_rf <- lcdfTrn %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
  mutate(predRet=(predict(rfModel_Ret, lcdfTrn))$predictions)
```

```
predRet_Trn_rf <- predRet_Trn_rf %>% mutate(tile=ntile(-predRet, 10))
```

```
predRet_Trn_rf %>% group_by(tile) %>% summarise(count=n(), avgpredRet=mean(predRet), numDefaults=
sum(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(act
ualTerm), totA=sum(grade=="A"), totB=sum(grade=="B"
), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgpredRet numDefaults avgActRet minRet maxRet avgTer totA
##   <int> <int>     <dbl>      <int>     <dbl>   <dbl>   <dbl> <dbl> <int>
## 1     1     1  5672     0.113         2    0.147  0.0905  0.367  0.995     6
## 2     2     2  5672     0.0881        13    0.110  0.0742  0.180  1.58    10
## 3     3     3  5672     0.0767        20    0.0932  0.0566  0.177  2.05    34
## 4     4     4  5672     0.0683        35    0.0815  0.0367  0.157  2.27   218
## 5     5     5  5672     0.0611        66    0.0725  0.0352  0.188  2.32   739
## 6     6     6  5671     0.0542        82    0.0644  0       0.121  2.26  1646
## 7     7     7  5671     0.0471       135    0.0539  0.0113  0.110  2.34  3021
## 8     8     8  5671     0.0402       177    0.0442 -0.00461 0.102  2.60  4498
## 9     9     9  5671     0.0226      2067    0.0138 -0.115   0.114  2.85  3566
## 10    10    10  5671    -0.0702     5671   -0.166 -0.333  -0.00326 3     545
## # ... with 5 more variables: totB <int>, totC <int>, totD <int>, totE <int>,
## #   totF <int>
```

*#Performance by deciles on Tst*

```
predRet_Tst_rf <- lcdfTst %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
  mutate(predRet=(predict(rfModel_Ret, lcdfTst))$predictions)
```

```
predRet_Tst_rf <- predRet_Tst_rf %>% mutate(tile=ntile(-predRet, 10))
```

```
predRet_Tst_rf %>% group_by(tile) %>% summarise(count=n(), avgpredRet=mean(predRet), numDefaults=
sum(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(act
ualTerm), totA=sum(grade=="A"), totB=sum(grade=="B"
), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgpredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>     <dbl>      <int>     <dbl>   <dbl>   <dbl> <dbl> <int> <int>
## 1     1     1  2431     0.0771        410    0.0747 -0.333  0.300  2.05     0    218
## 2     2     2  2431     0.0646        406    0.0622 -0.312  0.311  2.12     0    726
## 3     3     3  2431     0.0583        383    0.0587 -0.322  0.397  2.19    15    966
## 4     4     4  2431     0.0533        338    0.0566 -0.322  0.240  2.21   157   1088
## 5     5     5  2431     0.0489        348    0.0530 -0.309  0.248  2.20   497    977
## 6     6     6  2431     0.0450        304    0.0478 -0.322  0.279  2.24   860    807
## 7     7     7  2431     0.0413        346    0.0393 -0.323  0.265  2.29  1152    627
## 8     8     8  2430     0.0377        255    0.0448 -0.323  0.284  2.26  1335    497
## 9     9     9  2430     0.0327        321    0.0394 -0.313  0.271  2.33  1241    572
## 10    10    10  2430     0.0204        448    0.0380 -0.333  0.444  2.39   862    592
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

The ranger RF model performed well on Train Data, but never quite as good on Test Data. This is a sign of overfit. To avoid overfit, we tried several parameters to tune our ranger model: max.depth = (30, 20, 10, 6, 4, default of NULL), mtry (4, 6, 8), min.node.size = (10, 20, default 5), num.trees = (200, 500, 1000), sample.fraction = 0.5 and default of 1, replace=FALSE, respect.unordered.factors = "order", verbose = TRUE, seed=0. The problem with our experiment is that none of the changes of parameters lead to a change in the model's performance on test data.

When we included total\_pymnt in building the model, the performance increased significantly for Test Data. However, we are aware that total\_pymnt is a leakage variable, therefore we cannot include it in the model.

We created the decile table for Training and Test data using the ranger model. Our model is able to distinguish the loans with the highest actualReturn, which are concentrated on the top deciles. As we go down the deciles, the average predicted return decreases.

However, this model does not avoid defaults very well. Even though the Average Actual Return for the top deciles are the highest, there are too many defaults. Combining the Actual Return and Loan Status model is necessary.

```
library(xgboost)
```

```
##  
## Attaching package: 'xgboost'
```

```
## The following object is masked from 'package:dplyr':  
##  
## slice
```

```
set.seed(12)  
  
xDTrn_Ret <- lcdFTrn %>% select(-loan_status, -actualTerm, -annRet, -total_pymnt)  
xDTst_Ret <- lcdFTst %>% select(-loan_status, -actualTerm, -annRet, -total_pymnt)  
  
x = grep("actualReturn", colnames(xDTrn_Ret))  
  
train_x = data.matrix(xDTrn_Ret[, -x])  
train_y = data.matrix(xDTrn_Ret[, x])  
  
test_x = data.matrix(xDTst_Ret[, -x])  
test_y = data.matrix(xDTst_Ret[, x])  
  
test_500_x = test_x[1:500,]  
test_500_y = test_y[1:500,]  
  
xgb_train = xgb.DMatrix(data = train_x, label = train_y)  
xgb_test = xgb.DMatrix(data = test_x, label = test_y)  
xgb_500_test = xgb.DMatrix(data = test_500_x, label = test_500_y)  
  
xgbParam_1 <- list (max.depth = 8, eval_metric="error", eta = 0.1)  
  
#best model  
xgb_lsbest_ret <- xgb.train(xgbParam_1, xgb_train, nrounds = 1000)  
  
print(xgb_lsbest_ret)
```

```
## ##### xgb.Booster
## raw: 16.5 Mb
## call:
##   xgb.train(params = xgbParam_1, data = xgb_train, nrounds = 1000)
## params (as set within xgb.train):
##   max_depth = "8", eval_metric = "error", eta = "0.1", validate_parameters = "TRUE"
## xgb.attributes:
##   niter
## callbacks:
##   cb.print.evaluation(period = print_every_n)
## # of features: 41
## niter: 1000
## nfeatures : 41
```

```
pred_y = predict(xgb_lsbest_ret, xgb_test)


mse = mean((test_y - pred_y)^2)
mae = caret::MAE(test_y, pred_y)
rmse = caret::RMSE(test_y, pred_y)

cat("MSE: ", mse, "MAE: ", mae, " RMSE: ", rmse)
```

```
## MSE:  0.007521148 MAE:  0.05448671 RMSE:  0.08672455
```

```
x = 1:length(test_y)
plot(x, test_y, col = "red", type = "l", main = "Prediction on All Test Data Points")
lines(x, pred_y, col = "blue", type = "l")
legend(x = 1, y = 38, legend = c("original test_y", "predicted test_y"), col = c("red", "blue"),
box.lty = 1, cex = 0.8, lty = c(1, 1))
```


## Prediction on All Test Data Points



```
pred_500_y = predict(xgb_lsbest_ret, xgb_500_test)

x = 1:length(test_500_y)
plot(x, test_500_y, col = "red", type = "l", main = "Prediction on 500 Test Data Points")
lines(x, pred_500_y, col = "blue", type = "l")
legend(x = 1, y = 38, legend = c("original test_500_y", "predicted test_500_y"), col = c("red",
"blue"), box.lty = 1, cex = 0.8, lty = c(1, 1))
```

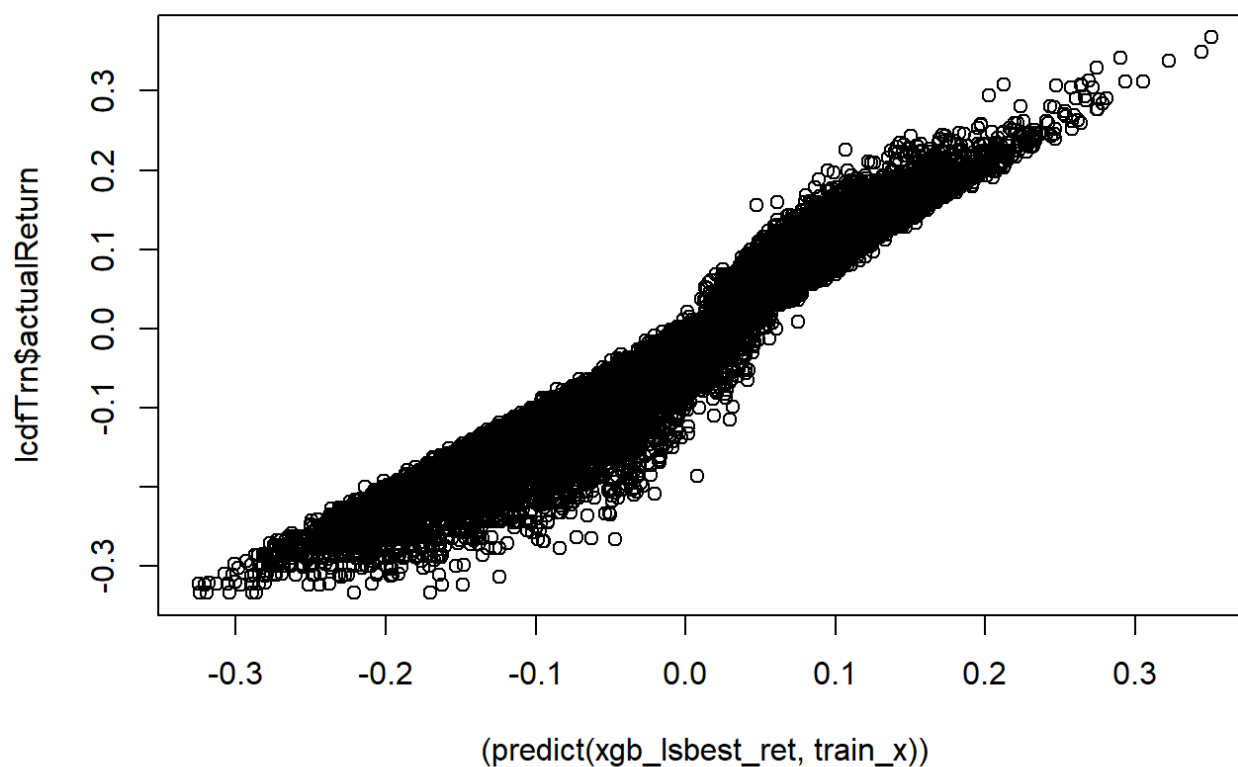
## Prediction on 500 Test Data Points



```
plot ( (predict(xgb_lsbest_ret, train_x)), lcdfTrn$actualReturn, main = "XGB Actual Return Model  
on Training Data")
```




## XGB Actual Return Model on Training Data



```
plot ( (predict(xgb_lsbest_ret, test_x)), lcdfTst$actualReturn, main = "XGB Actual Return Model on Test Data")
```

## XGB Actual Return Model on Test Data



```
#Performance by deciles on Trn
predRet_Trn_xgb <- lcdfTrn %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
mutate(predRet=(predict(xgb_lsbest_ret, train_x)))

predRet_Trn_xgb <- predRet_Trn_xgb %>% mutate(tile=ntile(-predRet, 10))

predRet_Trn_xgb %>% group_by(tile) %>% summarise(count=n(), avgpredRet=mean(predRet), numDefaults
=sum(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(act
ualTerm), totA=sum(grade=="A"), totB=sum(grade=="B"
), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgpredRet numDefaults avgActRet minRet maxRet avgTer totA
##   <int> <int>      <dbl>      <int>      <dbl>      <dbl>      <dbl> <dbl> <int>
## 1     1     1  5672      0.133         6    0.145    0.0809    0.367    1.19     3
## 2     2     2  5672      0.0988        12    0.108    0.0609    0.225    1.75    10
## 3     3     3  5672      0.0855        38    0.0932    0.0561    0.189    2.06    24
## 4     4     4  5672      0.0758        58    0.0819    0.00841   0.160    2.24   125
## 5     5     5  5672      0.0673        81    0.0731    0.0296    0.143    2.31   497
## 6     6     6  5671      0.0589       105    0.0645   -0.0123    0.160    2.22  1559
## 7     7     7  5671      0.0507       134    0.0542   -0.0141    0.155    2.28  3301
## 8     8     8  5671      0.0430       161    0.0453   -0.0652    0.100    2.46  4584
## 9     9     9  5671      0.0226      2002    0.0137   -0.209    0.0892    2.77  3662
## 10    10    10  5671     -0.120      5671   -0.164   -0.333   -0.0263     3    518
## # ... with 5 more variables: totB <int>, totC <int>, totD <int>, totE <int>,
## #   totF <int>
```

*#Performance by deciles on Tst*

```
predRet_Tst_xgb <- lcdfTst %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
mutate(predRet=(predict(xgb_lsbest_ret, test_x)))
```

```
predRet_Tst_xgb <- predRet_Tst_xgb %>% mutate(tile=ntile(-predRet, 10))
```

```
predRet_Tst_xgb %>% group_by(tile) %>% summarise(count=n(), avgpredRet=mean(predRet), numDefaults
=sum(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(act
ualTerm), totA=sum(grade=="A"), totB=sum(grade=="B"
), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgpredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>      <dbl>      <int>      <dbl>      <dbl>      <dbl> <dbl> <int> <int>
## 1     1     1  2431      0.0989       464    0.0665   -0.322    0.311    2.13     2   306
## 2     2     2  2431      0.0776       394    0.0611   -0.333    0.273    2.15    14   755
## 3     3     3  2431      0.0677       375    0.0584   -0.322    0.262    2.17    98   974
## 4     4     4  2431      0.0600       349    0.0545   -0.312    0.284    2.19   306  1012
## 5     5     5  2431      0.0532       311    0.0514   -0.323    0.265    2.21   661   920
## 6     6     6  2431      0.0470       290    0.0475   -0.308    0.309    2.23  1004   734
## 7     7     7  2431      0.0409       285    0.0450   -0.313    0.238    2.26  1153   641
## 8     8     8  2430      0.0339       292    0.0430   -0.312    0.271    2.28  1185   584
## 9     9     9  2430      0.0232       331    0.0454   -0.333    0.397    2.29   993   596
## 10    10    10  2430     -0.00624       468    0.0417   -0.323    0.444    2.37   703   548
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

source: <https://www.datatechnotes.com/2020/08/regression-example-with-xgboost-in-r.html>  
[\(https://www.datatechnotes.com/2020/08/regression-example-with-xgboost-in-r.html\)](https://www.datatechnotes.com/2020/08/regression-example-with-xgboost-in-r.html)

With the XGboost model, we tried the parameters max.depth and eta. eta - It makes the model more robust by shrinking the weights on each step. Typical values are 0.01-0.2 and the default value is 0.3. When experimenting with all the values and the parameters, we have troubles finding the correct combination to give better results on

Test Data. The change in parameters did not result in notable changes in performance for Test, but it did for Train Data. We proceed with  $\eta = 0.1$ ,  $n_{\text{rounds}} = 1000$ ,  $\text{max.depth} = 6$ . We did this because we needed to counter the high bias even in Training data predictions. This improved our performance in Train Data, but not Test Data (before this, the model was weak on both).

We created the decile table for Training and Test data using the XGB model. Our model is able to distinguish the loans with the highest actualReturn, which are concentrated on the top deciles. As we go down the deciles, the average predicted return decreases.

However, this model does not avoid defaults very well. Even though the Average Actual Return for the top deciles are the highest, there are too many defaults. Combining the Actual Return and Loan Status model is necessary.

## Q3: Combining Loan Status and Actual Return Model

Combining Loan Status and Actual Return models is necessary because: - Loan Status prediction alone only contains loans from the higher grades (A or B), resulting in low average actual return. This makes sense because loans from higher grades are safer from defaults, but have lower interest rates for investors. - Prediction from Actual Return alone only prioritizes the actual returns, but it does a poor job in avoiding defaults. Since the cost of a default could be much higher than any return, we should account for the defaults.

The combined predictions is made by: - Taking the prediction on actual return. - Adding the probability of a fully paid loan for each data point (score from loan status model). - Taking the top decile of this combination, meaning we are taking the loans those are predicted to have the highest returns. - Sorting the top decile by the score from loan status model, from highest probability of a fully paid to the lowest.

This best Loan Status models for RF and XGB were copied from last assignment's best model. Please refer to Q1 for best GLM model on Loan Status.

```
rgModel1 <- ranger(loan_status ~., data=subset(bs_lcdfTrn, select=-c(annRet, actualTerm, actualReturn, total_pymnt)), num.trees = 200, importance='permutation', mtry = 7, max.depth = 10, min.node.size = 30, sample.fraction = 0.5, replace=FALSE, respect.unordered.factors = "order" , verbose = TRUE , seed=0, probability = TRUE)
```

```
## Computing permutation importance.. Progress: 97%. Estimated remaining time: 1 seconds.
```

```
library(xgboost)
library(caret)

dumVar<-dummyVars(~.,data=lcdf %>% select(-loan_status))
dxlcdf<- predict(dumVar,lcdf)

# for loan_status, check levels and convert to dummy vars and keep the class label of interest
levels(lcdf$loan_status)
```

```
## [1] "Fully Paid" "Charged Off"
```

```

dylcdf <- class2ind(lcdf$loan_status, drop2nd = FALSE)
# and then decide which one to keep
colcdf <- dylcdf [ , 1]# or,fplcdf <- dyclcdf [ , 2]

#Training, test subsets
dxlcdfTrn <- dxlcdf[trnIndex,]
colcdfTrn <- colcdf[trnIndex]
dxlcdfTst <- dxlcdf[-trnIndex,]
colcdfTst <- colcdf[-trnIndex]
dxTrn <- xgb.DMatrix(subset(dxlcdfTrn, select=-c(annRet, actualTerm, actualReturn, total_pymnt)),
label=colcdfTrn)
dxTst <- xgb.DMatrix( subset( dxlcdfTst,select=-c(annRet, actualTerm, actualReturn, total_pymn
t)), label=colcdfTst)

#use cross-validation on training dataset to determine best model
xgbParam <- list (
max_depth = 4, eta = 0.01,
objective = "binary:logistic",
eval_metric="error", eval_metric = "auc")
xgb_lscv <- xgb.cv( xgbParam, dxTrn, nrounds = 10, nfold=10, early_stopping_rounds = 10 )

```

```

## [1] train-error:0.145766+0.000540   train-auc:0.679204+0.001121 test-error:0.145852+0.004941
test-auc:0.671531+0.011822
## Multiple eval metrics are present. Will use test_auc for early stopping.
## Will train until test_auc hasn't improved in 10 rounds.
##
## [2] train-error:0.145766+0.000540   train-auc:0.680188+0.001523 test-error:0.145834+0.004937
test-auc:0.671795+0.011558
## [3] train-error:0.145768+0.000536   train-auc:0.680967+0.001263 test-error:0.145834+0.004937
test-auc:0.672477+0.011810
## [4] train-error:0.145768+0.000536   train-auc:0.681378+0.001238 test-error:0.145834+0.004937
test-auc:0.672844+0.011531
## [5] train-error:0.145768+0.000536   train-auc:0.681748+0.001490 test-error:0.145834+0.004937
test-auc:0.673147+0.011112
## [6] train-error:0.145768+0.000536   train-auc:0.682545+0.001473 test-error:0.145834+0.004937
test-auc:0.673779+0.011615
## [7] train-error:0.145768+0.000536   train-auc:0.683019+0.001231 test-error:0.145834+0.004937
test-auc:0.674246+0.011497
## [8] train-error:0.145770+0.000531   train-auc:0.683275+0.001214 test-error:0.145799+0.004859
test-auc:0.674474+0.011431
## [9] train-error:0.145770+0.000531   train-auc:0.684062+0.001167 test-error:0.145799+0.004859
test-auc:0.675091+0.011158
## [10] train-error:0.145770+0.000531   train-auc:0.684271+0.001344 test-error:0.145799+0.004859
test-auc:0.675127+0.011018

```

```

#best iteration
xgb_lscv$best_iteration

```

```

## [1] 10

```

```
# or for the best iteration based on performance measure (among those specified in xgbParam)
best_cvIter <- which.max(xgb_lscv$evaluation_log$test_auc_mean)

#best model
xgb_lsbest <- xgb.train(xgbParam, dxTrn, nrounds = xgb_lscv$best_iteration)
```

## Combining predictions on ranger RF models

```
#Loan Status Model RF
```

```
rpredTst<-predict(rgModel1, lcdfTst)
scoreTst_FP <- rpredTst$predictions[, "Fully Paid"]

prPerfRF <- data.frame(scoreTst_FP)
prRetPerfRF <- cbind(prPerfRF, status=lcdfTst$loan_status, grade=lcdfTst$grade, actRet=lcdfTst$actualReturn, actTerm = lcdfTst$actualTerm)
prRetPerfRF <- prRetPerfRF %>% mutate(decile = ntile(-scoreTst_FP, 10))
prRetPerfRF %>% group_by(decile) %>% summarise(count=n(), numDefaults=sum(status=="Charged Off"),
avgActRet=mean(actRet),
minRet=min(actRet), maxRet=max(actRet), avgTer=mean(actTerm), totA=sum(grade=="A"), totB=sum(grade=="B" ), totC=sum(grade=="C"),
totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 13
##   decile count numDefaults avgActRet minRet maxRet avgTer totA totB totC
##   <int> <int>      <int>      <dbl> <dbl> <dbl> <dbl> <int> <int> <int>
## 1     1   2431         72    0.0408 -0.279  0.134  2.19  2411    20     0
## 2     2   2431        121    0.0432 -0.303  0.140  2.18  2050   377     4
## 3     3   2431        188    0.0490 -0.323  0.309  2.20  1079  1296    52
## 4     4   2431        229    0.0526 -0.322  0.238  2.22   443  1783   184
## 5     5   2431        297    0.0559 -0.333  0.218  2.20   118  1743   499
## 6     6   2431        351    0.0587 -0.312  0.217  2.20    18  1195  1083
## 7     7   2431        417    0.0605 -0.333  0.397  2.22     0   506  1639
## 8     8   2430        509    0.0550 -0.322  0.300  2.26     0   135  1706
## 9     9   2430        599    0.0517 -0.322  0.284  2.29     0    15  1314
## 10    10  2430        776    0.0472 -0.322  0.444  2.32     0     0   287
## # ... with 3 more variables: totD <int>, totE <int>, totF <int>
```

```
# Actual Return Model RF
```

```
predrfRet <- predict(rfModel_Ret, lcdfTst)
predrfRet_Tst <- lcdfTst %>% select(grade, loan_status, actualTerm, actualReturn, int_rate) %>%
mutate( predrfRet=predrfRet$predictions)
predrfRet_Tst <- predrfRet_Tst %>% mutate(tile=ntile(-predrfRet, 10))
predrfRet_Tst %>% group_by(tile) %>% summarise(count=n(), avgPredRet=mean(predrfRet),
numDefaults=sum(loan_status=="Charged Off"), avgActRet=mean(actualReturn), minRet=min(actualReturn),
maxRet=max(actualReturn), avgTer=mean(actualTerm), totA=sum(grade=="A"), totB=sum(grade=="B" ),
totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgPredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>     <dbl>         <int>     <dbl> <dbl> <dbl> <dbl> <int> <int>
## 1     1     1    2431     0.0771         410    0.0747 -0.333  0.300  2.05     0    218
## 2     2     2    2431     0.0646         406    0.0622 -0.312  0.311  2.12     0    726
## 3     3     3    2431     0.0583         383    0.0587 -0.322  0.397  2.19    15    966
## 4     4     4    2431     0.0533         338    0.0566 -0.322  0.240  2.21   157   1088
## 5     5     5    2431     0.0489         348    0.0530 -0.309  0.248  2.20   497    977
## 6     6     6    2431     0.0450         304    0.0478 -0.322  0.279  2.24   860    807
## 7     7     7    2431     0.0413         346    0.0393 -0.323  0.265  2.29  1152    627
## 8     8     8    2430     0.0377         255    0.0448 -0.323  0.284  2.26  1335    497
## 9     9     9    2430     0.0327         321    0.0394 -0.313  0.271  2.33  1241    572
## 10    10    2430     0.0204         448    0.0380 -0.333  0.444  2.39   862    592
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

```
#Combine both models using RF
```

```
d=1
pRetSc <- predrfRet_Tst %>% mutate(poScore=scoreTst_FP) #score scoreTst_rf_ls is from predicting
loan_status. predrfRet_Tst is predicting actual return
pRet_d <- pRetSc %>% filter(tile<=d)
pRet_d<- pRet_d %>% mutate(tile2=ntile(-poScore, 20))
pRet_d %>% group_by(tile2) %>% summarise(count=n(), avgPredRet=mean(predrfRet),
numDefaults=sum(loan_status=="Charged Off"), avgActRet=mean(actualReturn), minRet=min(actualReturn),
maxRet=max(actualReturn), avgTer=mean(actualTerm), totA=sum(grade=="A"), totB=sum(grade=="B" ),
totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 20 x 14
##   tile2 count avgPredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>    <dbl>         <int>    <dbl>    <dbl> <dbl> <dbl> <int> <int>
## 1     1     1    122     0.0747           9    0.0749 -0.234  0.194  1.96     0    80
## 2     2     2    122     0.0765           9    0.0867 -0.195  0.234  1.94     0    46
## 3     3     3    122     0.0769           6    0.0927 -0.125  0.217  1.99     0    33
## 4     4     4    122     0.0775          13    0.0764 -0.218  0.160  2.05     0    19
## 5     5     5    122     0.0766          12    0.0847 -0.233  0.198  1.81     0    17
## 6     6     6    122     0.0756          14    0.0757 -0.243  0.199  1.98     0    16
## 7     7     7    122     0.0770          17    0.0767 -0.277  0.182  2.04     0     4
## 8     8     8    122     0.0761          11    0.0921 -0.194  0.244  1.90     0     2
## 9     9     9    122     0.0761          18    0.0703 -0.277  0.244  2.09     0     1
## 10    10    10    122     0.0768          23    0.0609 -0.333  0.212  2.19     0     0
## 11    11    122     0.0778          24    0.0692 -0.253  0.198  2.22     0     0
## 12    12    121     0.0772          18    0.0849 -0.263  0.246  1.97     0     0
## 13    13    121     0.0768          22    0.0756 -0.218  0.277  2.08     0     0
## 14    14    121     0.0785          23    0.0761 -0.322  0.300  2.10     0     0
## 15    15    121     0.0769          26    0.0678 -0.275  0.251  2.11     0     0
## 16    16    121     0.0769          28    0.0680 -0.258  0.224  2.12     0     0
## 17    17    121     0.0785          28    0.0673 -0.251  0.262  2.05     0     0
## 18    18    121     0.0796          31    0.0750 -0.286  0.273  2.16     0     0
## 19    19    121     0.0778          38    0.0594 -0.310  0.256  2.06     0     0
## 20    20    121     0.0784          40    0.0592 -0.293  0.266  2.19     0     0
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

We first start by combining the predictions with Loan Status and Actual Returns models, and then we took the top decile from the combined models and created a new table from it. It turns out that the combined model performs better than the two original ones individually.

The model managed to differentiate 2431 loans from the lower grades (mostly C-E) with highest returns (averaged at 7%), much higher than that of the loan status model (4%). However, we need to avoid the abundance of defaults in this top decile from Actual Return prediction.

By further dividing this decile into 20 deciles, and then sorting them in decending order of their loan status prediction score, we can isolate the loans with the highest probability of being paid in the top deciles. The resulting table allows us to identify 122 loans with average predicted return at around 7.4% with only 10 defaults.

Although our model does not boast the highest possible accuracy, we can see that as we go down the deciles at the last table, the number of defaults go up significantly from 10 to 38 at the last deciles.

## Combining predictions on XGB Models



*#Predicting loan status using XGB*

```
xpredTst<-predict(xgb_lsbest, dxTst)

scoreTst_xgb_ls <- lcdfTst %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
mutate(score=xpredTst)

scoreTst_xgb_ls <- scoreTst_xgb_ls %>% mutate(tile=ntile(-score, 10))

scoreTst_xgb_ls %>% group_by(tile) %>% summarise(count=n(), avgSc=mean(score), numDefaults=sum(lo
an_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(act
ualTerm), totA=sum(grade=="A"),
totB=sum(grade=="B" ), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum
(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgSc numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int> <dbl>         <int>      <dbl> <dbl> <dbl> <dbl> <int> <int>
## 1     1     1  2431 0.544             97    0.0383 -0.313  0.115  2.22  2431     0
## 2     2     2  2431 0.542             118   0.0429 -0.313  0.142  2.19  2107    321
## 3     3     3  2431 0.540             196   0.0445 -0.323  0.175  2.25  1307   1087
## 4     4     4  2431 0.538             262   0.0536 -0.312  0.309  2.19   274   1914
## 5     5     5  2431 0.536             295   0.0570 -0.333  0.238  2.28     0   1795
## 6     6     6  2431 0.533             375   0.0516 -0.323  0.248  2.17     0   1800
## 7     7     7  2431 0.531             441   0.0622 -0.322  0.277  2.13     0    134
## 8     8     8  2430 0.529             496   0.0567 -0.322  0.397  2.26     0     0
## 9     9     9  2430 0.527             538   0.0548 -0.333  0.284  2.28     0     1
## 10    10    10  2430 0.519             741   0.0529 -0.322  0.444  2.32     0    18
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

*#Predicting Actual Return using XGB*

```
xpredTst_ret<-predict(xgb_lsbest_ret, xgb_test)
predXgbRet_Tst <- lcdfTst %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
mutate( predXgbRet=xpredTst_ret)
predXgbRet_Tst <- predXgbRet_Tst %>% mutate(tile=ntile(-predXgbRet, 10))
predXgbRet_Tst %>% group_by(tile) %>% summarise(count=n(), avgPredRet=mean(predXgbRet),
numDefaults=sum(loan_status=="Charged Off"), avgActRet=mean(actualReturn), minRet=min(actualRetu
rn),
maxRet=max(actualReturn), avgTer=mean(actualTerm), totA=sum(grade=="A"), totB=sum(grade=="B" ),
totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgPredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>      <dbl>      <int>      <dbl>  <dbl>  <dbl>  <dbl> <int> <int>
## 1     1     1  2431    0.0989        464    0.0665 -0.322  0.311  2.13     2   306
## 2     2     2  2431    0.0776        394    0.0611 -0.333  0.273  2.15    14   755
## 3     3     3  2431    0.0677        375    0.0584 -0.322  0.262  2.17    98   974
## 4     4     4  2431    0.0600        349    0.0545 -0.312  0.284  2.19   306  1012
## 5     5     5  2431    0.0532        311    0.0514 -0.323  0.265  2.21   661   920
## 6     6     6  2431    0.0470        290    0.0475 -0.308  0.309  2.23  1004   734
## 7     7     7  2431    0.0409        285    0.0450 -0.313  0.238  2.26  1153   641
## 8     8     8  2430    0.0339        292    0.0430 -0.312  0.271  2.28  1185   584
## 9     9     9  2430    0.0232        331    0.0454 -0.333  0.397  2.29   993   596
## 10    10    10  2430   -0.00624        468    0.0417 -0.323  0.444  2.37   703   548
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

*#Combining both models*

```
d=1
pRetSc_xgb <- predXgbRet_Tst %>% mutate(poScore=scoreTst_xgb_ls$score)
#score scoreTst_xgb_ls is from predicting loan_status. predXgbRet_Tst is predicting actual return

pRet_d_xgb <- pRetSc_xgb %>% filter(tile<=d)
pRet_d_xgb <- pRet_d_xgb %>% mutate(tile2=ntile(-poScore, 20))
pRet_d_xgb %>% group_by(tile2) %>% summarise(count=n(), avgPredRet=mean(xpredTst_ret),
numDefaults=sum(loan_status=="Charged Off"), avgActRet=mean(actualReturn), minRet=min(actualReturn),
maxRet=max(actualReturn), avgTer=mean(actualTerm), totA=sum(grade=="A"), totB=sum(grade=="B" ),
totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 20 x 14
##   tile2 count avgPredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>   <dbl>         <int>   <dbl>   <dbl> <dbl> <dbl> <int> <int>
## 1     1    122    0.0496          11    0.0620 -0.300  0.126  2.14     2    95
## 2     2    122    0.0496          10    0.0688 -0.222  0.187  1.97     0   103
## 3     3    122    0.0496          11    0.0839 -0.205  0.234  1.98     0    19
## 4     4    122    0.0496          14    0.0847 -0.194  0.244  2.01     0    20
## 5     5    122    0.0496          14    0.0708 -0.230  0.185  1.94     0    68
## 6     6    122    0.0496          29    0.0555 -0.277  0.212  1.92     0     0
## 7     7    122    0.0496          16    0.0753 -0.266  0.187  2.04     0     0
## 8     8    122    0.0496          16    0.0768 -0.282  0.225  2.00     0     0
## 9     9    122    0.0496          24    0.0693 -0.293  0.225  2.09     0     0
## 10    10    122    0.0496          12    0.0845 -0.215  0.239  2.21     0     0
## 11    11    122    0.0496          25    0.0619 -0.264  0.281  2.12     0     0
## 12    12    121    0.0496          27    0.0535 -0.285  0.210  2.26     0     0
## 13    13    121    0.0496          22    0.0570 -0.276  0.159  2.35     0     0
## 14    14    121    0.0496          19    0.0704 -0.265  0.190  2.18     0     0
## 15    15    121    0.0496          31    0.0617 -0.297  0.271  2.29     0     0
## 16    16    121    0.0496          37    0.0498 -0.273  0.212  2.19     0     0
## 17    17    121    0.0496          35    0.0519 -0.322  0.230  2.18     0     0
## 18    18    121    0.0496          36    0.0591 -0.285  0.311  2.20     0     1
## 19    19    121    0.0496          33    0.0748 -0.321  0.253  2.20     0     0
## 20    20    121    0.0496          42    0.0569 -0.297  0.300  2.26     0     0
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

We followed the same method to combine the predictions of both XGB models.

The combined XGB models managed to differentiate loans from the lower grades with minimum defaults. The combined models did better than the individual models: - The Loan Status model alone was able to minimize defaults on the top deciles but only has 3.8% average actual return. - The Actual Return model alone managed to gather an average actual return of 6.6% (much higher than the Loan Status prediction), in the expense of an overwhelming amount of defaults (464 loans).

We further dissected this top decile into 20 new deciles, and sort them by their probability of being fully paid (score from the loan status model).

The top 122 loans only has 11 defaults with average actual return of 6.2%. The combined deciles allowed us to focus on the top 122 loans with the minimum defaults which present an investment opportunity with comparatively lower risk with good actual returns that are better than what we saw in individual models. As we go down the 20 deciles, we can see that the number of defaults gradually goes up to 42 out of 122 loans. This means that our model has predictive capabilities, despite its shortcomings.

## Combining predictions on GLM Models

```
# Loan Status Decile with glm
xDTst<-lcdfTst %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_pymnt)

glmPredls_Tst=predict(glm1sw_cv,data.matrix(xDTst), s="lambda.min", type="response" )

preds_glm_Tst <- prediction(glmPredls_Tst, lcdfTst$loan_status, label.ordering = c("Charged Off",
"Fully Paid"))

scoreTst_glm_ls <- lcdfTst %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
mutate(score=glmPredls_Tst)

scoreTst_glm_ls <- scoreTst_glm_ls %>% mutate(tile=ntile(-score, 10))

scoreTst_glm_ls %>% group_by(tile) %>% summarise(count=n(), avgpredRet=mean(score), numDefaults=s
um(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(act
ualTerm), totA=sum(grade=="A"),
totB=sum(grade=="B" ), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum
(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgpredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>   <dbl>         <int>   <dbl>  <dbl>  <dbl>  <dbl> <int> <int>
## 1     1     1  2431     0.809           79  0.0419 -0.323  0.154  2.17 2218  199
## 2     2     2  2431     0.728          114  0.0463 -0.281  0.200  2.20 1855  537
## 3     3     3  2431     0.672          182  0.0473 -0.313  0.309  2.23 1252 1068
## 4     4     4  2431     0.619          245  0.0512 -0.333  0.238  2.19  564 1534
## 5     5     5  2431     0.567          295  0.0560 -0.322  0.239  2.18  174 1541
## 6     6     6  2431     0.519          356  0.0572 -0.322  0.244  2.21   47 1160
## 7     7     7  2431     0.471          436  0.0572 -0.333  0.271  2.22    6  658
## 8     8     8  2430     0.419          495  0.0554 -0.323  0.277  2.26    2  285
## 9     9     9  2430     0.357          571  0.0544 -0.322  0.397  2.29    1   86
## 10    10    10  2430     0.254          786  0.0476 -0.321  0.444  2.34    0    2
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

```
# Actual Return Decile with glm
glmPred_Tst_ret <- predict(glmRet_cv, data.matrix(lcdfTst %>% select(-loan_status, -actualTerm, -
annRet, -actualReturn, -total_pymnt)),s="lambda.min")
predRet_Tst_glm <- lcdfTst %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
mutate(predRet= glmPred_Tst_ret)

predRet_Tst_glm <- predRet_Tst_glm %>% mutate(tile=ntile(-predRet, 10))

predRet_Tst_glm %>% group_by(tile) %>% summarise(count=n(), avgpredRet=mean(predRet), numDefaults
=sum(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(act
ualTerm), totA=sum(grade=="A"),
totB=sum(grade=="B" ), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum
(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
```

```
##   tile count avgpredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>    <dbl>      <int>    <dbl>  <dbl>  <dbl>  <dbl> <int> <int>
## 1     1     1    2431    0.0733      449    0.0733 -0.322  0.444  2.08    14   345
## 2     2     2    2431    0.0635      396    0.0667 -0.333  0.334  2.13    71   583
## 3     3     3    2431    0.0589      386    0.0601 -0.333  0.277  2.14   190   712
## 4     4     4    2431    0.0555      393    0.0538 -0.322  0.397  2.19   310   785
## 5     5     5    2431    0.0524      358    0.0516 -0.312  0.262  2.23   441   805
## 6     6     6    2431    0.0495      374    0.0462 -0.309  0.223  2.24   604   815
## 7     7     7    2431    0.0467      330    0.0465 -0.312  0.248  2.27   805   783
## 8     8     8    2430    0.0436      311    0.0416 -0.312  0.238  2.31   993   788
## 9     9     9    2430    0.0397      273    0.0408 -0.322  0.208  2.31  1204   749
## 10    10    2430    0.0322      289    0.0341 -0.323  0.230  2.37  1487   705
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

```
d=1
```

```
pRetSc_glm <- predRet_Tst_glm %>% mutate(poScore=scoreTst_glm_ls$score) #score scoreTst_glm_ls is
from predicting loan_status. predRet_Tst_glm is predicting actual return
```

```
pRet_d_glm <- pRetSc_glm %>% filter(tile<=d)
```

```
pRet_d_glm <- pRet_d_glm %>% mutate(tile2=ntile(-poScore, 20))
```

```
pRet_d_glm %>% group_by(tile2) %>% summarise(count=n(), avgPredRet=mean(glmPred_Tst_ret),
numDefaults=sum(loan_status=="Charged Off"), avgActRet=mean(actualReturn), minRet=min(actualReturn),
maxRet=max(actualReturn), avgTer=mean(actualTerm), totA=sum(grade=="A"), totB=sum(grade=="B" ),
totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 20 x 14
##   tile2 count avgPredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>    <dbl>         <int>    <dbl>    <dbl> <dbl> <dbl> <int> <int>
## 1     1    122    0.0515           6    0.0717 -0.210  0.177  1.93    14    80
## 2     2    122    0.0515           7    0.0796 -0.234  0.200  1.91     0    79
## 3     3    122    0.0515           9    0.0699 -0.223  0.152  2.12     0    76
## 4     4    122    0.0515          13    0.0726 -0.276  0.229  2.04     0    60
## 5     5    122    0.0515          14    0.0742 -0.234  0.234  2.00     0    27
## 6     6    122    0.0515          14    0.0830 -0.265  0.207  1.87     0    15
## 7     7    122    0.0515          11    0.0836 -0.322  0.239  1.98     0     6
## 8     8    122    0.0515          23    0.0650 -0.287  0.223  2.12     0     2
## 9     9    122    0.0515          15    0.0835 -0.223  0.188  1.98     0     0
## 10    10    122    0.0515          20    0.0777 -0.230  0.185  2.25     0     0
## 11    11    122    0.0515          30    0.0529 -0.298  0.199  2.13     0     0
## 12    12    121    0.0515          23    0.0762 -0.297  0.271  2.01     0     0
## 13    13    121    0.0515          20    0.0929 -0.297  0.244  2.03     0     0
## 14    14    121    0.0515          24    0.0740 -0.284  0.268  2.15     0     0
## 15    15    121    0.0515          30    0.0743 -0.298  0.246  2.16     0     0
## 16    16    121    0.0515          32    0.0692 -0.293  0.245  2.14     0     0
## 17    17    121    0.0515          33    0.0657 -0.285  0.272  2.20     0     0
## 18    18    121    0.0515          32    0.0816 -0.321  0.272  2.01     0     0
## 19    19    121    0.0515          44    0.0561 -0.284  0.311  2.29     0     0
## 20    20    121    0.0515          49    0.0628 -0.320  0.444  2.31     0     0
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

We pretty much saw similar trend in case of GLM as what we saw in xgb and Ranger. Model where loan status was target variable we saw low number of defaults(79) but relatively very low actual return(4.2%) whereas the model with Actual return as target variable had higher average actual return of 7.3 percent but it at a price of higher number of defaults (449) which is actually in line with our expectations.

Following the methods that we applied for RF and XGboost, we now combine our GLM models. Below is the results of our individual models: - the top decile of the loan status prediction managed to identify the safest loans, with only 79 defaults out of 2431. However, the average actual return is so only 4.2% - the prediction from our Actual Return model managed to predict the highest average actual return of 7.3%, while failing to separate the defaults from non-defaults (default of 449 loans).

We partitioned the top decile of the Actual Return prediction into 20 new deciles. We also added the score of loan status prediction and sort the new 20 deciles based on this score. The resulting combination decile managed to identify the 2 top deciles (122 loans in each) with 7.1% and 7.9% average actual return with only 6 and 7 defaults each. So far, the glm model has been the best performer for combination predictions.

## Q4: Modelling with lower grade loans

Another approach to get the best returns from fully paid loans is by building a model to predict loan status on the lower grade loans. We will compare the results with the combination predictions above.

## Modelling from lower loan grades with Ranger

```
lg_lcdfTst<-lcdfTst %>% filter(grade=='C'| grade=='D'| grade== 'E'| grade== 'F'| grade== 'G')
lg_lcdfTrn<-lcdfTrn %>% filter(grade=='C'| grade=='D'| grade== 'E'| grade== 'F'| grade== 'G')

rf_M1_lg <- ranger(loan_status ~., data=subset(lg_lcdfTrn, select=-c(annRet, actualTerm, actualReturn)), num.trees =200,
probability=TRUE, importance='permutation')
lg_scoreTstRF <- lg_lcdfTst %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>%
mutate(score=(predict(rf_M1_lg,lg_lcdfTst))$predictions[, "Fully Paid"])

lg_scoreTstRF <- lg_scoreTstRF %>% mutate(tile=ntile(-score, 10))

lg_scoreTstRF %>% group_by(tile) %>% summarise(count=n(), avgSc=mean(score),
numDefaults=sum(loan_status=="Charged Off"), avgActRet=mean(actualReturn), minRet=min(actualReturn),
maxRet=max(actualReturn), avgTer=mean(actualTerm), totA=sum(grade=="A"), totB=sum(grade=="B" ),
totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgSc numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int> <dbl>      <int>      <dbl>    <dbl>  <dbl>  <dbl> <int> <int>
## 1     1     1  1112 0.960         16    0.0888 1.52e-2 0.217    2.35     0     0
## 2     2     2  1112 0.935         15    0.0926 1.95e-2 0.185    2.28     0     0
## 3     3     3  1112 0.917         23    0.0973 -3.58e-5 0.239    2.15     0     0
## 4     4     4  1112 0.899         39    0.0994 -5.13e-2 0.234    2.18     0     0
## 5     5     5  1112 0.879         44    0.104  -5.28e-2 0.277    2.04     0     0
## 6     6     6  1112 0.856         52    0.106  -1.05e-1 0.281    2.02     0     0
## 7     7     7  1112 0.826         80    0.110  -7.38e-2 0.256    1.96     0     0
## 8     8     8  1112 0.769        144    0.115  -1.52e-1 0.397    1.75     0     0
## 9     9     9  1111 0.526        900   -0.0478 -2.89e-1 0.444    2.64     0     0
## 10    10    1111 0.190       1111   -0.189  -3.33e-1 -0.0533    3         0     0
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

Our RF model on lower grade loans is able to identify the loans which are most likely to be fully paid (15 defaults out of 1112 loans) while maintaining a high average actual return of 8.9%.

```

xTrn_lg<-lg_lcdfTrn %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_pymnt)
yTrn_lg<-factor(if_else(lg_lcdfTrn$loan_status=="Fully Paid", '1', '0'))
wts_lg <- if_else(yTrn_lg == 0, 1-sum(yTrn_lg == 0)/length(yTrn_lg), 1-sum(yTrn_lg == 1)/length(yTrn_lg))

glm1sw_cv_lg<- cv.glmnet(data.matrix(xTrn_lg), yTrn_lg, family= "binomial", weights = wts_lg, alpha = 1)

xTst_lg<-lg_lcdfTst %>% select(-loan_status, -actualTerm, -annRet, -actualReturn, -total_pymnt)

glmPredls_Tst_lg=predict(glm1sw_cv_lg,data.matrix(xTst_lg), s="lambda.min", type="response" )

preds_glm_Tst_lg <- prediction(glmPredls_Tst_lg, lg_lcdfTst$loan_status, label.ordering = c("Charged Off", "Fully Paid"))

scoreTst_glm_ls_lg <- lg_lcdfTst %>% select(grade, loan_status, actualReturn, actualTerm, interest_rate) %>% mutate(score=glmPredls_Tst_lg)

scoreTst_glm_ls_lg <- scoreTst_glm_ls_lg %>% mutate(tile=ntile(-score, 10))

scoreTst_glm_ls_lg %>% group_by(tile) %>% summarise(count=n(), avgpredRet=mean(score), numDefaults=sum(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(actualTerm), totA=sum(grade=="A"),
totB=sum(grade=="B" ), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F" ) )

```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```

## # A tibble: 10 x 14
##   tile count avgpredRet numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int>     <dbl>      <int>     <dbl>  <dbl>  <dbl>  <dbl> <int> <int>
## 1     1     1  1112      0.677        125  0.0729 -0.333  0.239  2.07     0     0
## 2     2     2  1112      0.614        184  0.0617 -0.322  0.244  2.16     0     0
## 3     3     3  1112      0.582        177  0.0659 -0.333  0.271  2.16     0     0
## 4     4     4  1112      0.555        197  0.0621 -0.322  0.208  2.21     0     0
## 5     5     5  1112      0.531        225  0.0581 -0.322  0.271  2.23     0     0
## 6     6     6  1112      0.506        255  0.0544 -0.322  0.397  2.24     0     0
## 7     7     7  1112      0.480        245  0.0571 -0.298  0.277  2.26     0     0
## 8     8     8  1112      0.448        286  0.0530 -0.322  0.284  2.32     0     0
## 9     9     9  1111      0.409        331  0.0470 -0.310  0.296  2.33     0     0
## 10    10    1111      0.328        399  0.0445 -0.321  0.444  2.37     0     0
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>

```

Our GLM model on lower grade loans is doing less compared to the our RF model on lower grade loans. However, its top decile has a good average actual return, which is 7.3% with only 125 defaults out of 1112 loans.



```
lg_lcdf<-lcdf %>% filter(grade=='C'| grade=='D'| grade== 'E'| grade== 'F'| grade== 'G')

nr_lg=nrow(lg_lcdf)
trnIndex_lg = sample(1:nr_lg, size = round(0.7*nr_lg), replace=FALSE)
lcdfTrn_lg=lg_lcdf[trnIndex_lg,]
lcdfTst_lg = lg_lcdf[-trnIndex_lg,]

dumVar_lg<-dummyVars(~.,data=lg_lcdf %>% select(-loan_status))
dxlcdf_lg<- predict(dumVar_lg,lg_lcdf)

# for loan_status, check levels and convert to dummy vars and keep the class label of interest
levels(lg_lcdf$loan_status)
```

```
## [1] "Fully Paid" "Charged Off"
```

```
dylcdf_lg <- class2ind(lg_lcdf$loan_status, drop2nd = FALSE)
# and then decide which one to keep
colcdf_lg <- dylcdf_lg [ , 1]# or, fplcdf <- dycldf [ , 2]

#Training, test subsets
dxlcdfTrn_lg <- dxlcdf_lg[trnIndex_lg,]
colcdfTrn_lg <- colcdf_lg[trnIndex_lg]
dxlcdfTst_lg <- dxlcdf_lg[-trnIndex_lg,]
colcdfTst_lg <- colcdf_lg[-trnIndex_lg]
dxTrn_lg <- xgb.DMatrix(subset(dxlcdfTrn_lg, select=-c(annRet, actualTerm, actualReturn, total_pymnt)), label=colcdfTrn_lg)
dxTst_lg <- xgb.DMatrix( subset( dxlcdfTst_lg,select=-c(annRet, actualTerm, actualReturn, total_pymnt)), label=colcdfTst_lg)

#use cross-validation on training dataset to determine best model
xgbParam_lg <- list (
max_depth = 4, eta = 0.01,
objective = "binary:logistic",
eval_metric="error", eval_metric = "auc")
xgb_lscv_lg <- xgb.cv( xgbParam_lg, dxTrn_lg, nrounds = 10, nfold=10, early_stopping_rounds = 10
)
```

```
## [1] train-error:0.214359+0.000709 train-auc:0.605757+0.002406 test-error:0.214815+0.005985
test-auc:0.583065+0.011516
## Multiple eval metrics are present. Will use test_auc for early stopping.
## Will train until test_auc hasn't improved in 10 rounds.
##
## [2] train-error:0.214380+0.000688 train-auc:0.606434+0.003375 test-error:0.214738+0.005934
test-auc:0.583626+0.011221
## [3] train-error:0.214388+0.000678 train-auc:0.608528+0.004824 test-error:0.214777+0.005956
test-auc:0.585515+0.012265
## [4] train-error:0.214380+0.000688 train-auc:0.609129+0.004719 test-error:0.214738+0.005934
test-auc:0.586059+0.011969
## [5] train-error:0.214380+0.000688 train-auc:0.609929+0.004977 test-error:0.214738+0.005934
test-auc:0.586818+0.012446
## [6] train-error:0.214393+0.000672 train-auc:0.611182+0.004533 test-error:0.214623+0.005789
test-auc:0.588042+0.012515
## [7] train-error:0.214393+0.000672 train-auc:0.612531+0.003918 test-error:0.214623+0.005789
test-auc:0.589305+0.012696
## [8] train-error:0.214401+0.000681 train-auc:0.613162+0.003844 test-error:0.214623+0.005789
test-auc:0.589874+0.012443
## [9] train-error:0.214414+0.000694 train-auc:0.614026+0.003905 test-error:0.214623+0.005789
test-auc:0.590258+0.012631
## [10] train-error:0.214431+0.000702 train-auc:0.614627+0.003259 test-error:0.214546+0.005753
test-auc:0.590537+0.013903
```

```
#best iteration
xgb_lscv_lg$best_iteration
```

```
## [1] 10
```

```
# or for the best iteration based on performance measure (among those specified in xgbParam)
best_cvIter_lg <- which.max(xgb_lscv_lg$evaluation_log$test_auc_mean)

#best model
xgb_lsbest_lg <- xgb.train(xgbParam_lg, dxTrn_lg, nrounds = xgb_lscv_lg$best_iteration)

xpredTst_lg<-predict(xgb_lsbest_lg, dxTst_lg)

scoreTst_xgb_ls_lg <- lcdfTst_lg %>% select(grade, loan_status, actualReturn, actualTerm, int_rate) %>% mutate(score=xpredTst_lg)

scoreTst_xgb_ls_lg <- scoreTst_xgb_ls_lg %>% mutate(tile=ntile(-score, 10))

scoreTst_xgb_ls_lg %>% group_by(tile) %>% summarise(count=n(), avgSc=mean(score), numDefaults=sum(loan_status=="Charged Off"),
avgActRet=mean(actualReturn), minRet=min(actualReturn), maxRet=max(actualReturn), avgTer=mean(actualTerm), totA=sum(grade=="A"),
totB=sum(grade=="B" ), totC=sum(grade=="C"), totD=sum(grade=="D"), totE=sum(grade=="E"), totF=sum(grade=="F") )
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 14
##   tile count avgSc numDefaults avgActRet minRet maxRet avgTer totA totB
##   <int> <int> <dbl>         <int>      <dbl> <dbl> <dbl> <dbl> <int> <int>
## 1     1    1 1117 0.535           159    0.0709 -0.310  0.294  2.11     0     0
## 2     2    2 1117 0.533           174    0.0634 -0.300  0.239  2.11     0     0
## 3     3    3 1117 0.532           180    0.0606 -0.300  0.217  2.24     0     0
## 4     4    4 1117 0.530           223    0.0536 -0.333  0.214  2.19     0     0
## 5     5    5 1117 0.529           215    0.0557 -0.293  0.244  2.18     0     0
## 6     6    6 1117 0.528           265    0.0533 -0.322  0.214  2.25     0     0
## 7     7    7 1116 0.527           256    0.0548 -0.322  0.268  2.27     0     0
## 8     8    8 1116 0.522           287    0.0539 -0.320  0.300  2.32     0     0
## 9     9    9 1116 0.521           311    0.0504 -0.333  0.290  2.29     0     0
## 10    10   10 1116 0.516           377    0.0523 -0.321  0.444  2.32     0     0
## # ... with 4 more variables: totC <int>, totD <int>, totE <int>, totF <int>
```

Our XGBoost model on lower grade loans is doing less compared to the our RF and GLM models on lower grade loans. However, its top decile has a decent average actual return of 6.6% with only 144 defaults out of 1112 loans. As we go down the deciles, we encounter more defaults.

## In comparing RF to RF, GLM to GLM, XGB to XGB:

1. Ranger Model seemed be working exceptionally well lower grade loans data. If we compare the top deciles of the two different Ranger models, one with lower grades loans and the other that contains all the grades: The Lower grade loan data model seems to be working better in all aspects be it default rate or the average actual return. Investing in lower grades loans seems to be a better choice if we go by the ranger model.
2. Comparison of xgb model with lower grade loan data with its counterpart that contains loans from all grade is pretty much in line with our expectations. Lower grade loan model has higher defaults but also higher actual average returns.
3. Comparison of glm model with lower grade loan data with its counterpart that contains loans from all grade is pretty much in line with our expectations. Lower grade loan model has higher defaults but also higher actual average returns.

## Profit analysis for best investment approach

Now we will apply all the models through the profit analysis. Suppose we have to decide which top decile is the best to invest on (using what model?), and we are planning to invest \$100 in every loan. The cost of investing in a Charged Off loan is \$35, as was explained in our previous assignment. The return of a loan will follow the percentage of average actual return for that decile.

We would recommend the following models to use, because of their best profit calculation: 1. Use the RF model on lower grade loans, which boasts the highest profit. 2. Use the combined GLM models, and invest in the top 122 loans.

The complete tabulation for investment approaches using the best models can be found in the table in the Appendix: Best investment approach based on earned profit.

Appendix- Variable Importance for different models

Variables are ranked pretty much similarly in most of the models with GLM being an exception. Top 5 variables are common in Ranger and XGBoost followed by little bit of variation in variable rank after that. GLMNet does have Grade in top 5 as one of the most important variables. But rest of the variables in the GLMNet list are different from other two models. Proportion of Loan Amount to Annual Income is apparently important in GLMNet model, this trend is clearly different from what we saw in other models. The effect of Multicollinearity can be clearly seen in GLMNet model as variable grade did not come out to be as important as we expected it to be.

Ranger	Importance factor	Importance NORM
int_rate	0.05	100%
sub_grade	0.04	89%
grade	0.03	60%
dti	0.02	46%
avg_cur_bal	0.02	34%
tot_hi_cred_lim	0.01	31%
acc_open_past_24mths	0.01	25%
annual_inc	0.01	24%
total_bc_limit	0.01	23%
bc_open_to_buy	0.01	23%
installment	0.01	23%
total_rev_hi_lim	0.01	17%
loan_amnt	0.01	14%
bc_util	0.01	14%
mort_acc	0.01	13%
mths_since_recent_inq	0.01	11%
total_bal_ex_mort	0.01	11%
mo_sin_old_rev_tl_op	0.01	11%
emp_length	0.00	11%
mo_sin_rcnt_tl	0.00	11%
num_op_rev_tl	0.00	10%
mths_since_recent_bc	0.00	10%
prop_OpAccts_to_TotAccts	0.00	10%
total_il_high_credit_limit	0.00	10%
home_ownership	0.00	10%
mo_sin_old_il_acct	0.00	9%
num_sats	0.00	8%
num_rev_accts	0.00	8%
num_bc_sats	0.00	7%
mo_sin_rcnt_rev_tl_op	0.00	6%
num_bc_tl	0.00	6%
num_il_tl	0.00	6%
propSatisBankcardAccts	0.00	5%
pct_tl_nvr_dlq	0.00	5%
purpose	0.00	3%
initial_list_status	0.00	1%
tax_liens	0.00	0%
chargeoff_within_12_mths	0.00	0%
delinq_amnt	0.00	0%
num_tl_30dpd	0.00	0%
num_tl_30dpd	0.00	0%

XGBoost	Importance	Importance NORM
int_rate	0.8010646	100%
grade	0.09476281	12%
dti	0.03740318	5%
avg_cur_bal	0.02048943	3%
tot_hi_cred_lim	0.01702559	2%
propLoanAmt_to_AnnInc	0.0165261	2%
acc_open_past_24mths	0.00614935	1%
emp_length.n/a	0.00313037	0%
installment	0.00168584	0%
chargeoff_within_12_mths	0.00114687	0%
annual_inc	0.00061587	0%

--	--	--

GLMNet	Importance	Importance NORM
propLoanAmt_to_AnnInc	61%	100%
num_tl_30dpd	26%	42%
prop_OpAccts_to_TotAccts	19%	32%
grade	11%	18%
home_ownership	7%	11%
propSatisBankcardAccts	7%	11%
int_rate	6%	9%
acc_open_past_24mths	5%	8%
mort_acc	3%	5%
num_bc_tl	3%	4%
num_op_rev_tl	3%	4%
tax_liens	2%	4%
sub_grade	2%	4%
dti	2%	3%
num_rev_accts	1%	2%
initial_list_status	1%	2%
chargeoff_within_12_mths	1%	2%
emp_length	1%	2%
num_bc_sats	1%	1%
num_sats	0%	1%
purpose	0%	1%
pct_tl_nvr_dlq	0%	1%
bc_util	0%	0%
num_il_tl	0%	0%
mths_since_recent_inq	0%	0%
mo_sin_rcnt_tl	0%	0%
mths_since_recent_bc	0%	0%
mo_sin_old_rev_tl_op	0%	0%
installment	0%	0%
mo_sin_old_il_acct	0%	0%
delinq_amnt	0%	0%
mo_sin_rcnt_rev_tl_op	0%	0%
total_bc_limit	0%	0%
bc_open_to_buy	0%	0%
total_il_high_credit_limit	0%	0%
total_rev_hi_lim	0%	0%
total_bal_ex_mort	0%	0%
avg_cur_bal	0%	0%
annual_inc	0%	0%
tot_hi_cred_lim	0%	0%

Best investment approach based on earned profit.

Models	nLoans	defaults	avgActualReturn	Investment	Loss_from_CO	Profit_from_FP	Profit/Loss
Combined RF	122	10	0.07	12200	-350	832.94	482.94
Combined XGB	122	25	0.06	12200	-875	600.33	-274.67
Combined GLM	122	6	0.07	12200	-210	831.95	621.95
Lower Grade RF	1112	15	0.09	111200	-525	9684.04	9159.04
Lower Grade XGB	1112	144	0.07	111200	-5040	6395.12	1355.12
Lower Grade GLM	1112	125	0.07	111200	-4375	7195.63	2820.63